

Intro

Human vs AI

# Atari Games: Human VS AI

The leaderboard: <https://github.com/cshenton/atari-leaderboard>


Game	Top Human Score	Top Machine Score	Best	Best Machine	Learning Type	Notes
Alien	103583	9491	Human	Rainbow	Q-gradient	
Amidar	71529	5131	Human	Rainbow	Q-gradient	
Assault	8647	14497	Machine	A3C	Policy-gradient	
Asterix	1000000	428200	Human	Rainbow	Q-gradient	
Asteroids	57340	5093	Human	A3C	Policy-gradient	*
Atlantis	10604840	2311815	Human	PPO	Policy-gradient	
Bank Heist	45899	1611	Human	Dueling DDQN	Q-gradient	
Battlezone	98000	62010	Human	Rainbow	Q-gradient	
Beamrider	52866	26172	Human	Prioritized DDQN	Q-gradient	1B
Berzerk	1057940	2545	Human	Rainbow	Q-gradient	
Bowling	279	135	Human	HyperNEAT	Genetic Policy	J

# Section 1

## Gym Environment

# Gym Env

...implements the classic “agent-environment loop”



# Gym Env


...is a standard API for reinforcement learning...

```
import gym
env = gym.make("LunarLander-v2", render_mode="human")
observation, info = env.reset(seed=42)
for _ in range(1000):
    action = policy(observation) # User-defined policy function
    observation, reward, terminated, truncated, info = env.step(action)

    if terminated or truncated:
        observation, info = env.reset()
env.close()
```

# Gym Env


...is a standard API for reinforcement learning...






# Gym Env

...and a diverse collection of reference environments




[Adventure](#)




[Air Raid](#)




[Alien](#)




[Ant](#)



[Half Cheetah](#)




[Hopper](#)




[Amidar](#)




[Assault](#)




[Asterix](#)



[Humanoid Standup](#)



[Humanoid](#)




[Inverted Double Pendulum](#)




[Asteroids](#)




[Atlantis](#)




[Bank Heist](#)



[Inverted Pendulum](#)



[Reacher](#)



[Swimmer](#)

# Gym Env

## Demo-1




# Section 2

## Ray RLlib

# RAY RLlib

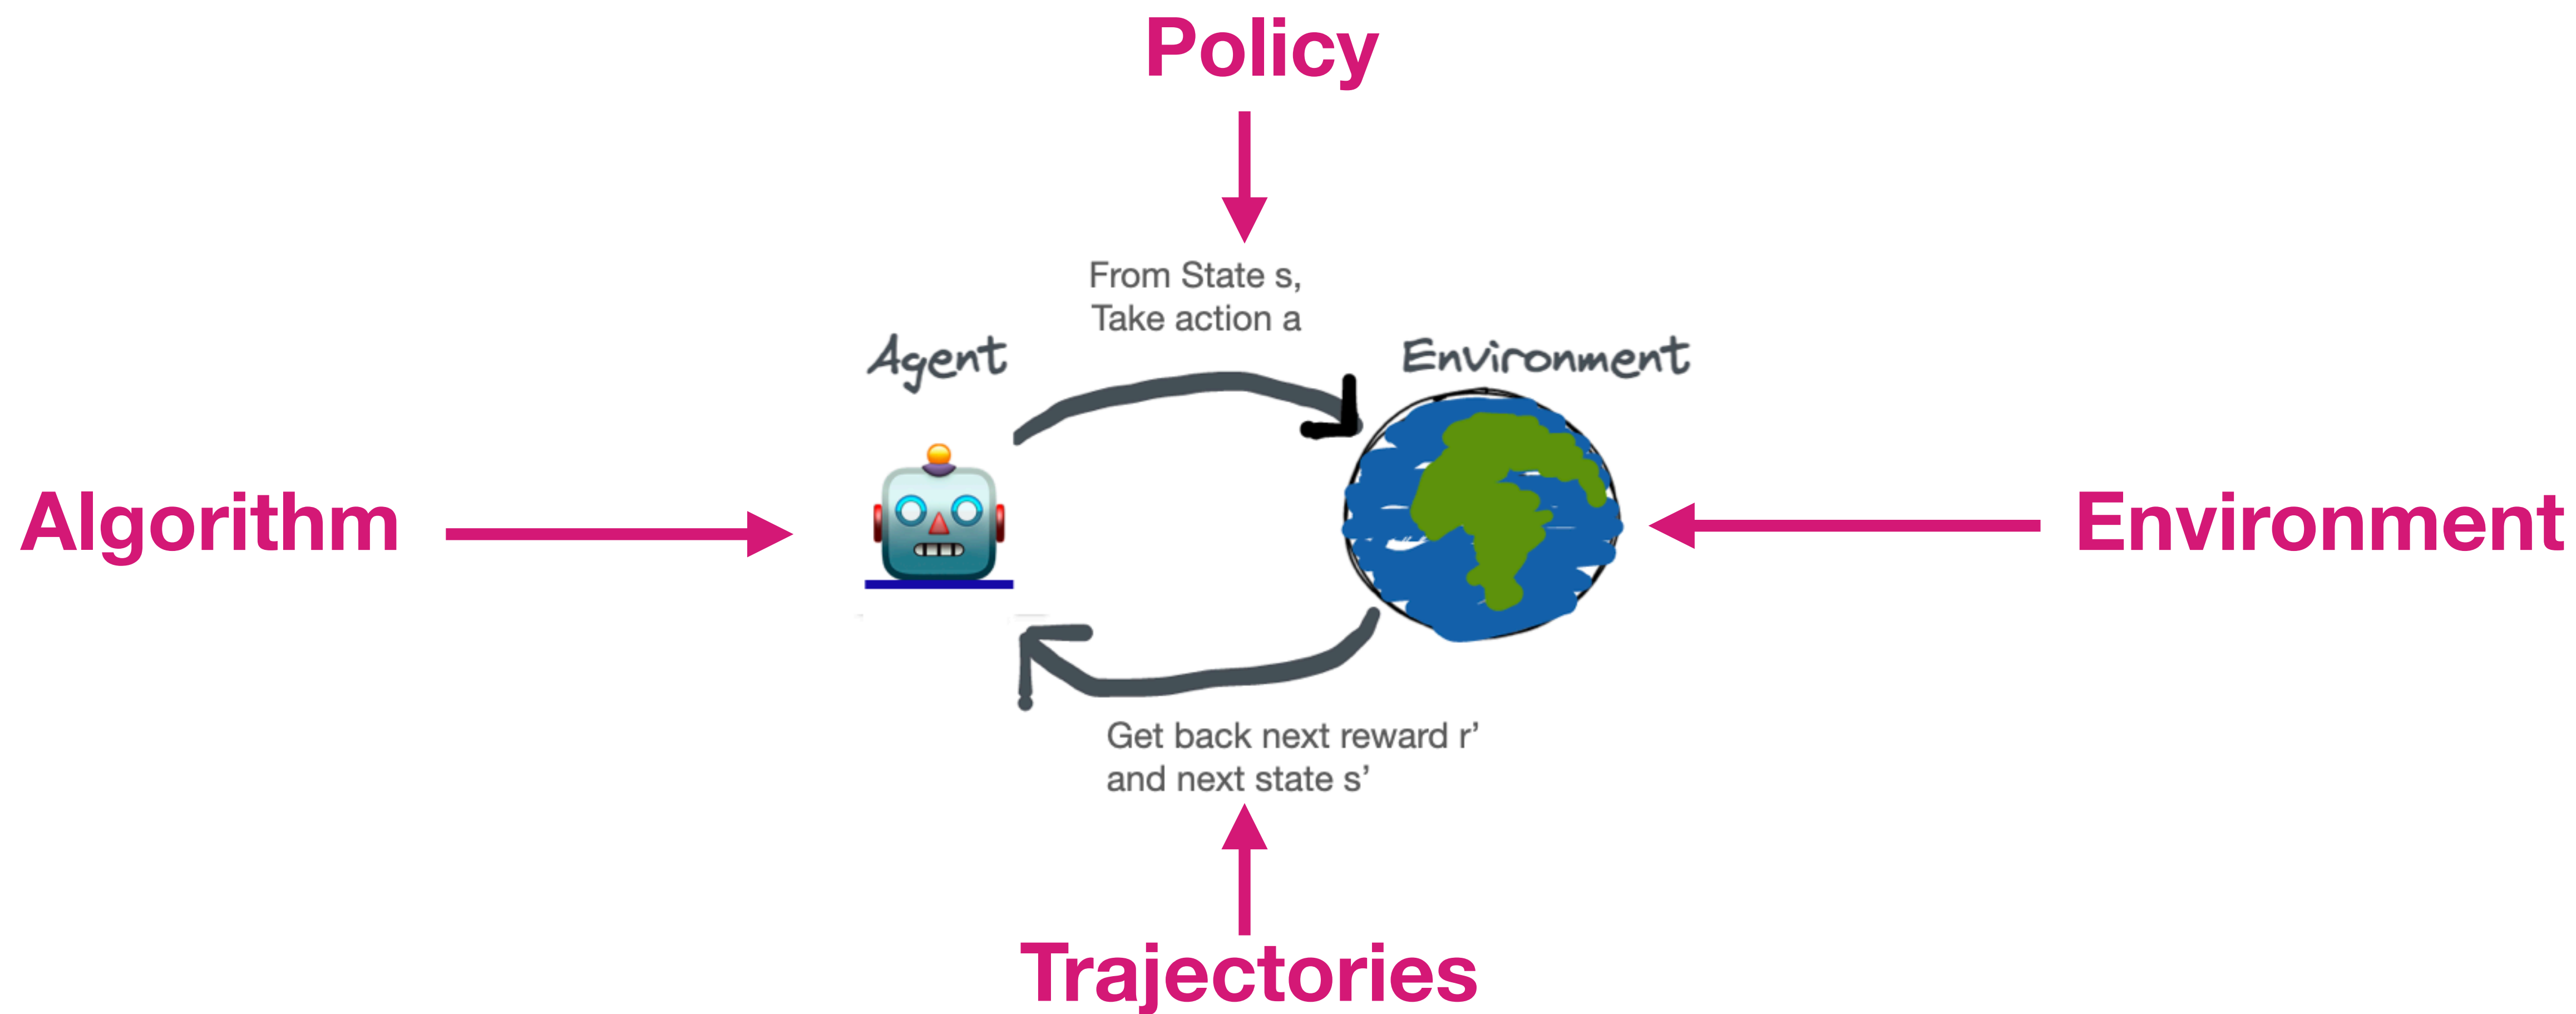
Industry-Grade Reinforcement Learning



# RAY

# RAY RLlib

## Components



# RAY RLlib

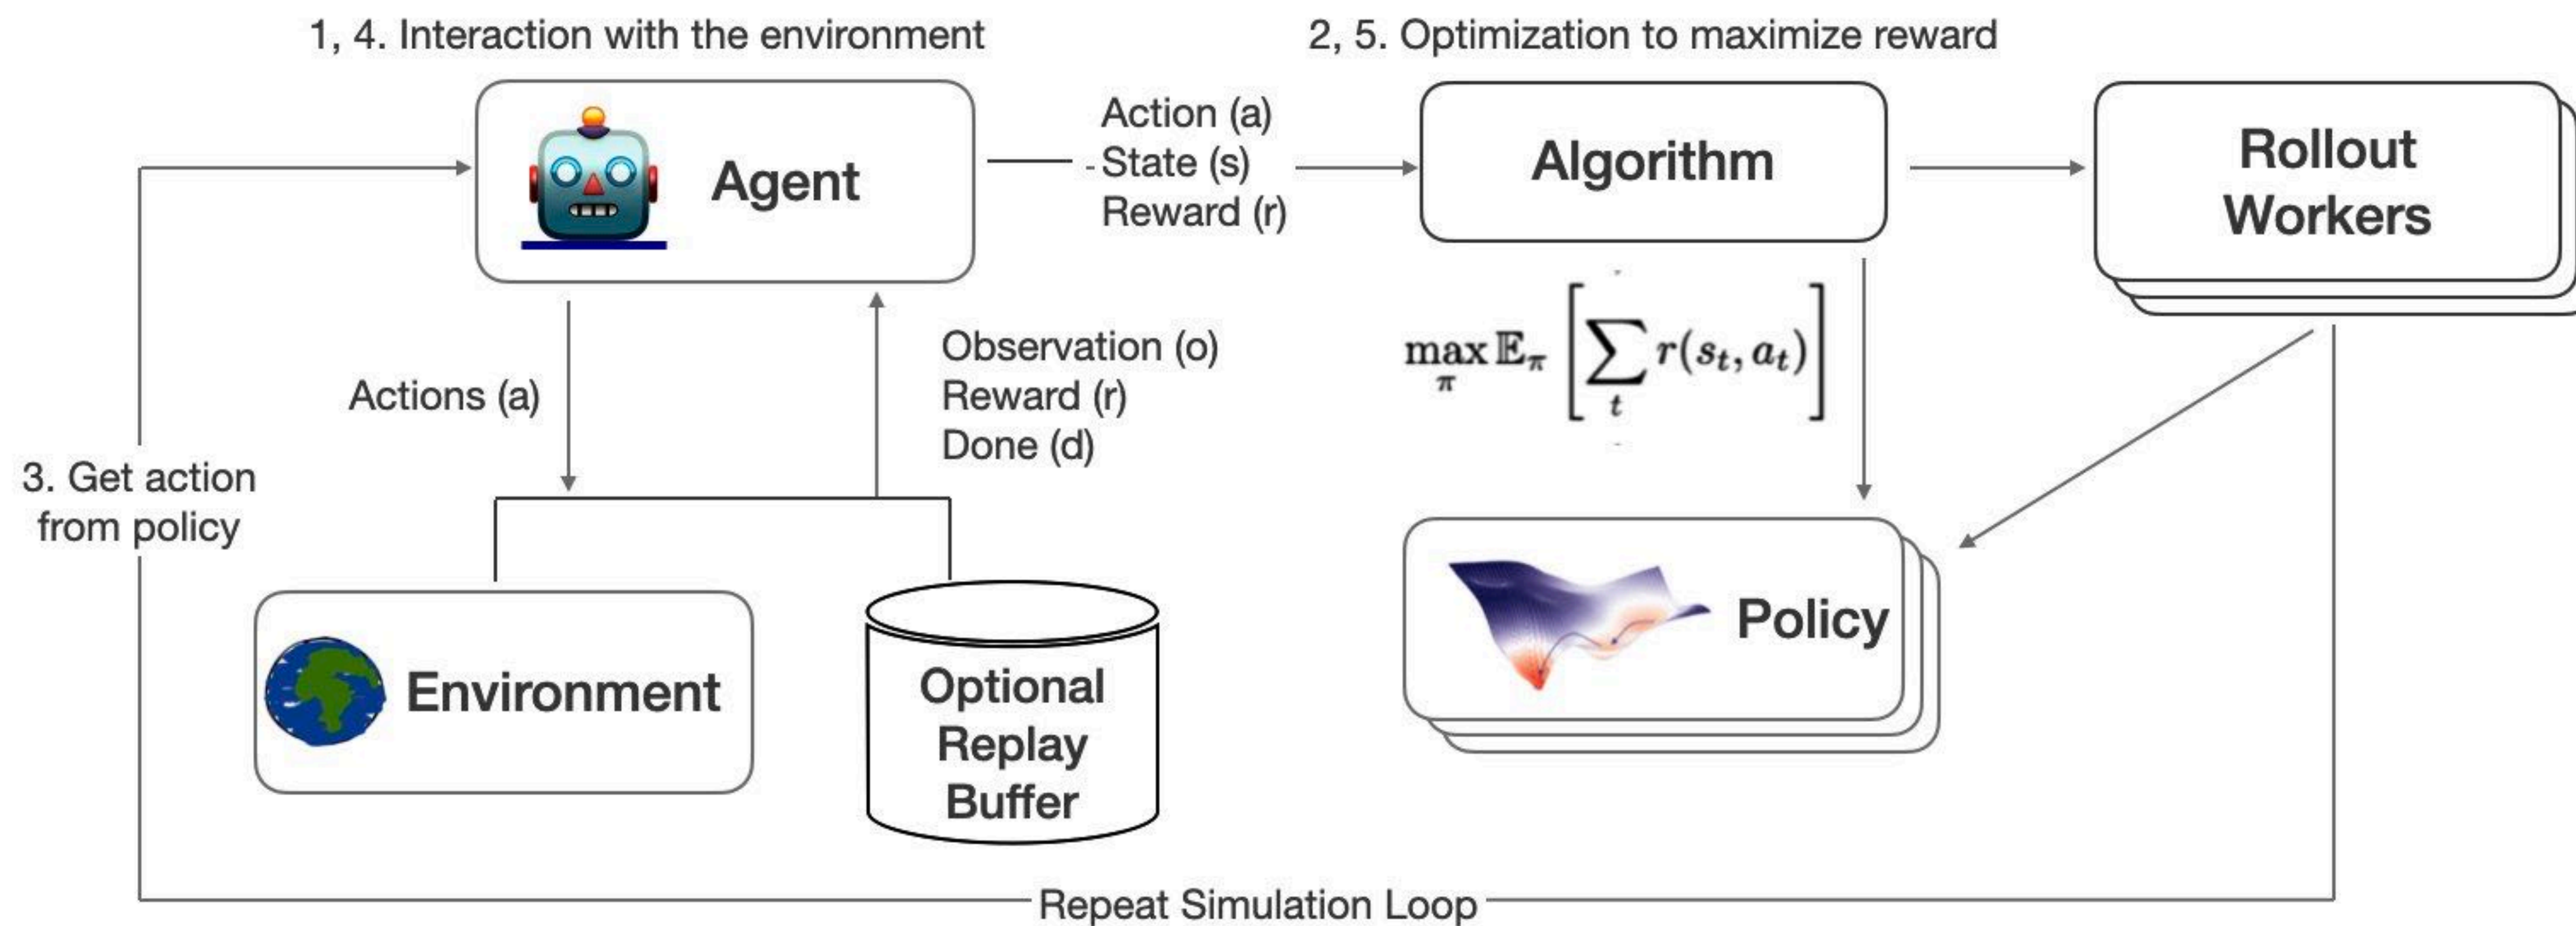
## Environment

- All possible actions (**action space**)
- **A** complete description of the environment, nothing hidden (**state space**)
- An observation by the agent of certain parts of the state (**observation space**)
- **Reward**, which is the only feedback the agent receives per action

The model that tries to maximize the expected sum over all future rewards is called a **policy**. The policy is a function mapping the environment's observations to an action to take, usually written  $\pi(s(t)) \rightarrow a(t)$ .

# RAY RLlib

## A diagram of the RL iterative learning process





# RAY RLlib

## Industry-Grade Reinforcement Learning



```
# Import the RL algorithm (Algorithm) we would like to use.
from ray.rllib.algorithms.ppo import PPO

# Configure the algorithm.
config = {
    # Environment (RLlib understands openAI gym registered strings).
    "env": "Taxi-v3",
    # Use 2 environment workers (aka "rollout workers") that parallelly
    # collect samples from their own environment clone(s).
    "num_workers": 2,
    # Change this to "framework: torch", if you are using PyTorch.
    # Also, use "framework: tf2" for tf2.x eager execution.
    "framework": "tf",
    # Tweak the default model provided automatically by RLlib,
    # given the environment's observation- and action spaces.
    "model": {
        "fcnet_hiddens": [64, 64],
        "fcnet_activation": "relu",
    },
    # Set up a separate evaluation worker set for the
    # `algo.evaluate()` call after training (see below).
    "evaluation_num_workers": 1,
    # Only for evaluation runs, render the env.
    "evaluation_config": {
        "render_env": True,
    },
}

# Create our RLlib Trainer.
algo = PPO(config=config)

# Run it for n training iterations. A training iteration includes
# parallel sample collection by the environment workers as well as
# loss calculation on the collected batch and a model update.
for _ in range(3):
    print(algo.train())

# Evaluate the trained Trainer (and render each timestep to the shell's
# output).
algo.evaluate()
```

# Ray RLlib

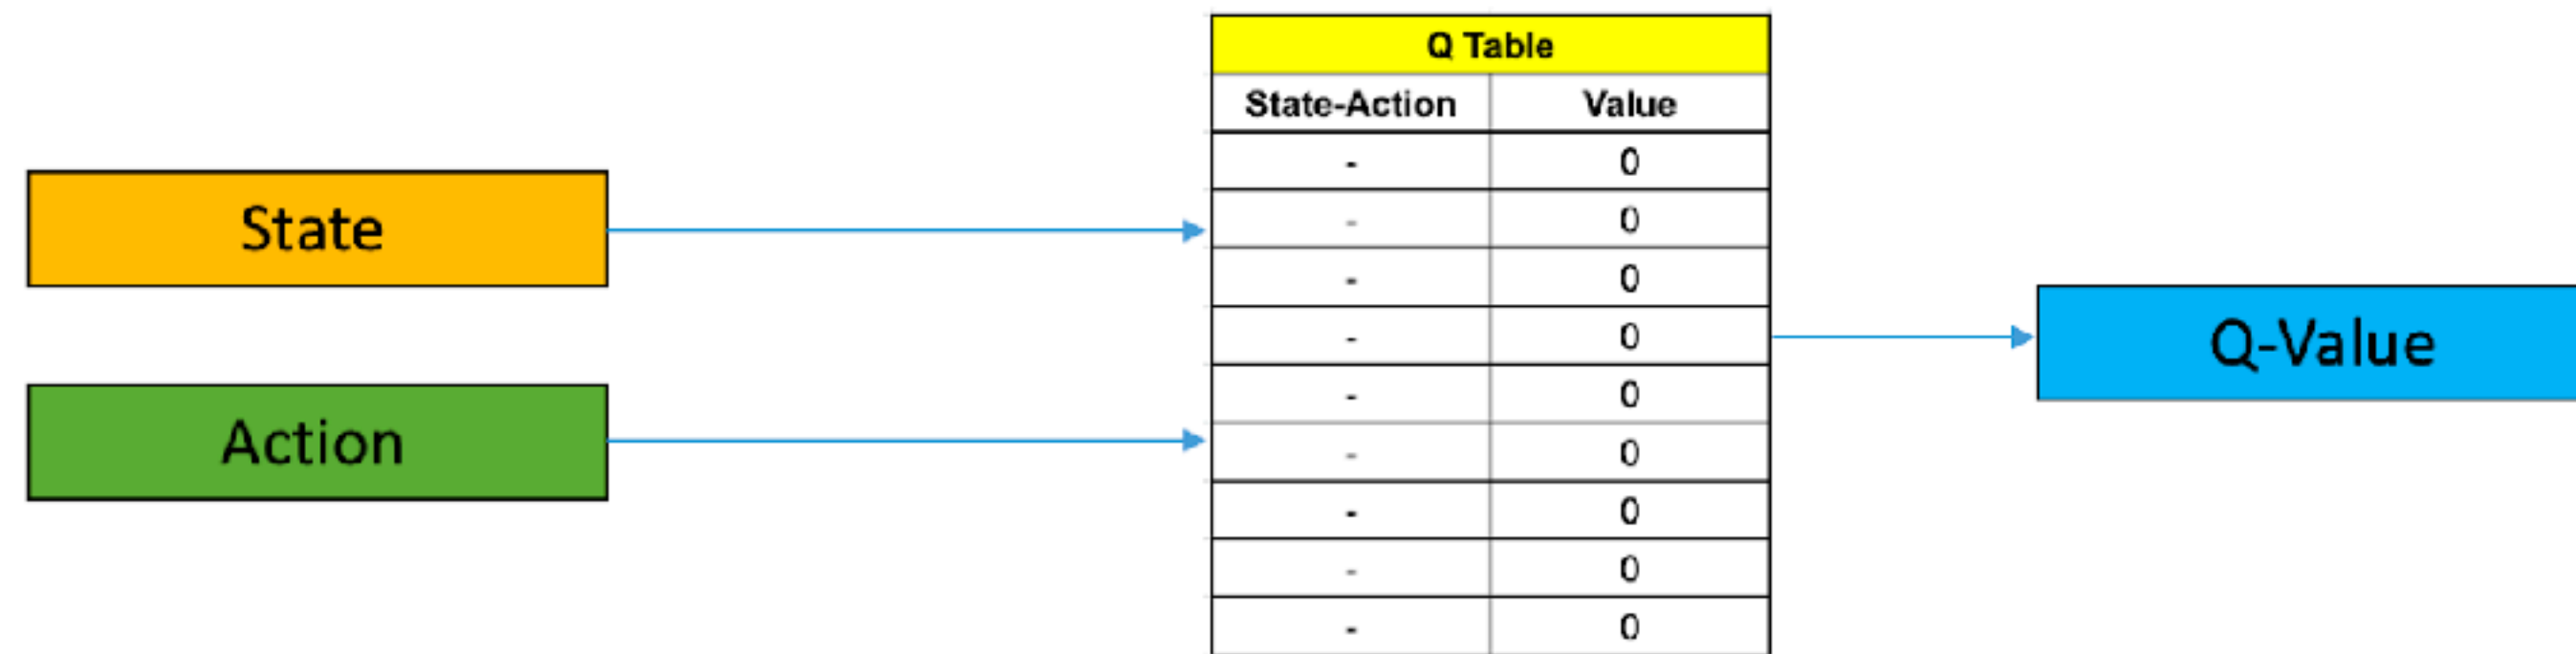
## Demo-2

# Section 3


DQN, Double DQN, APEX-DQN,  
Rainbow

# DQN = Deep Q-Network (1st gen)

Predicts Q-values (expected future rewards) for state-action pairs



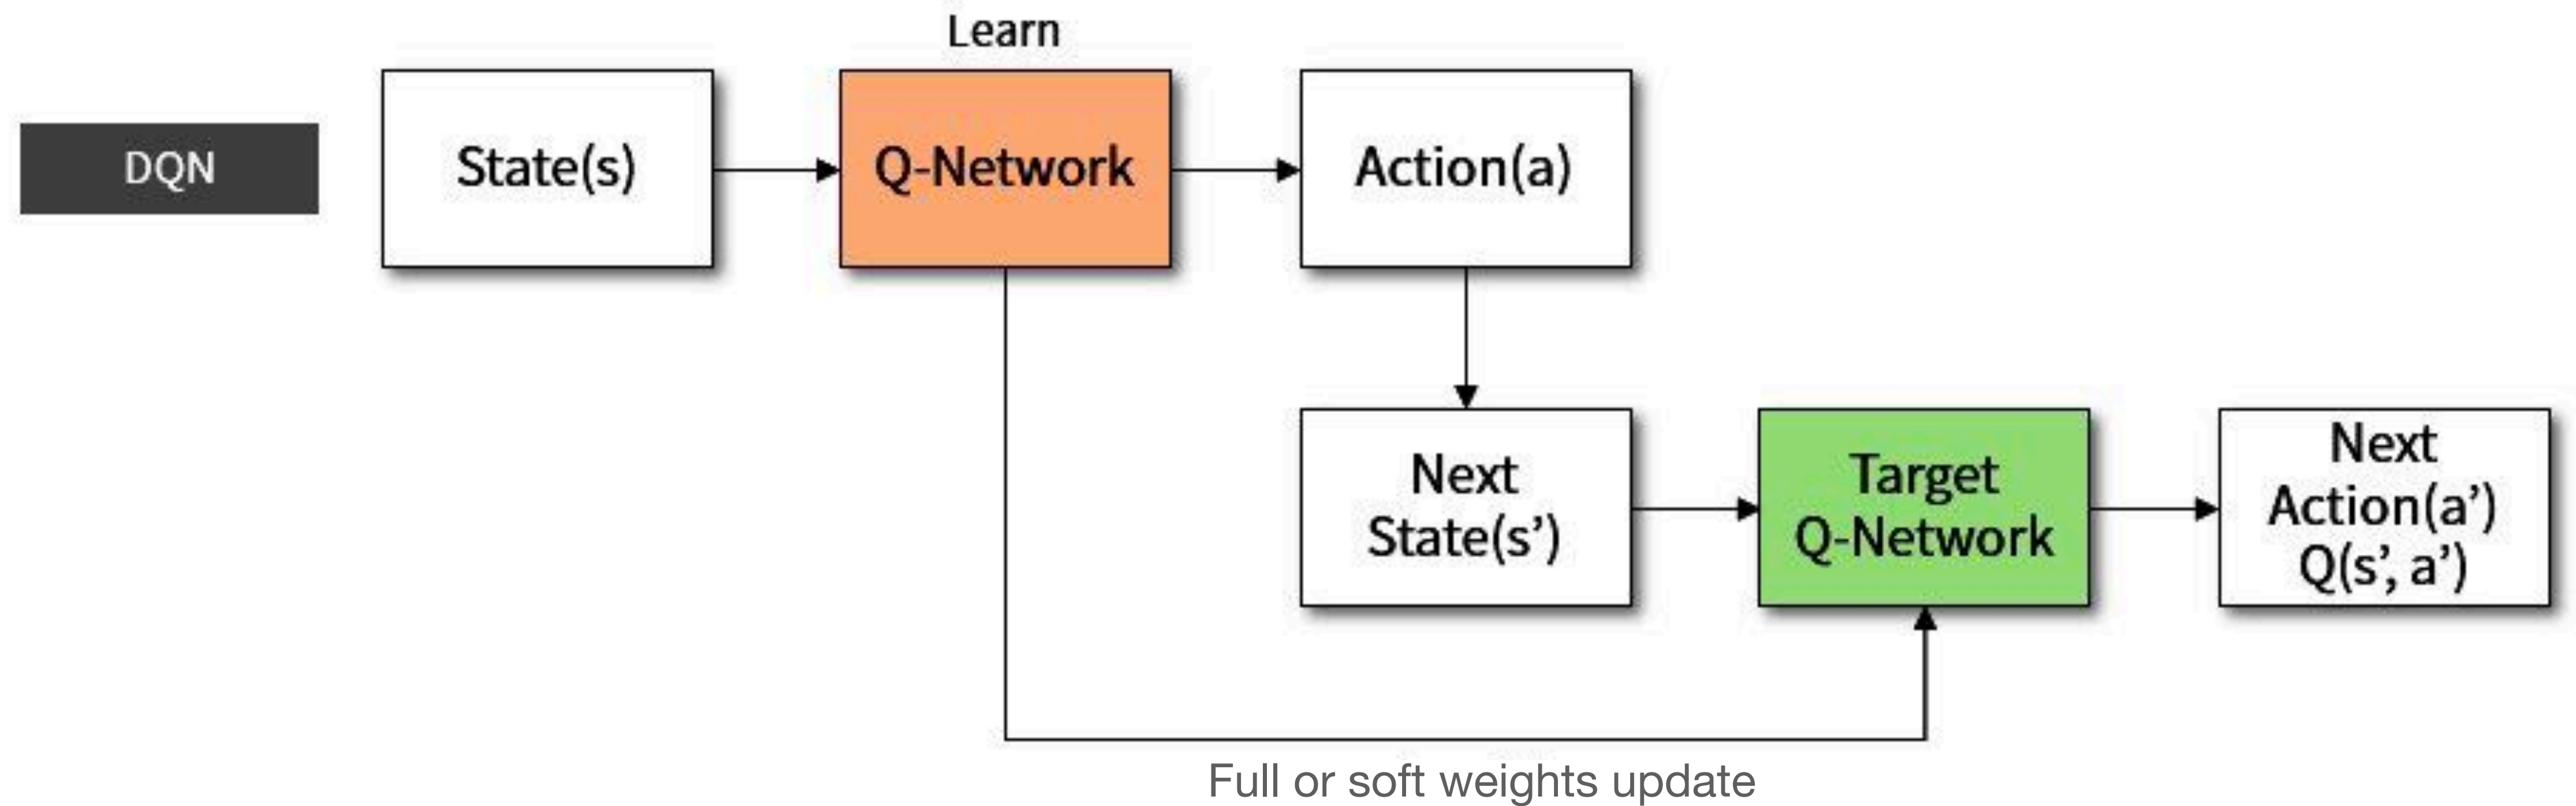
Q Learning



Deep Q Learning

# DQN = Deep Q-Network

## Training (one network)



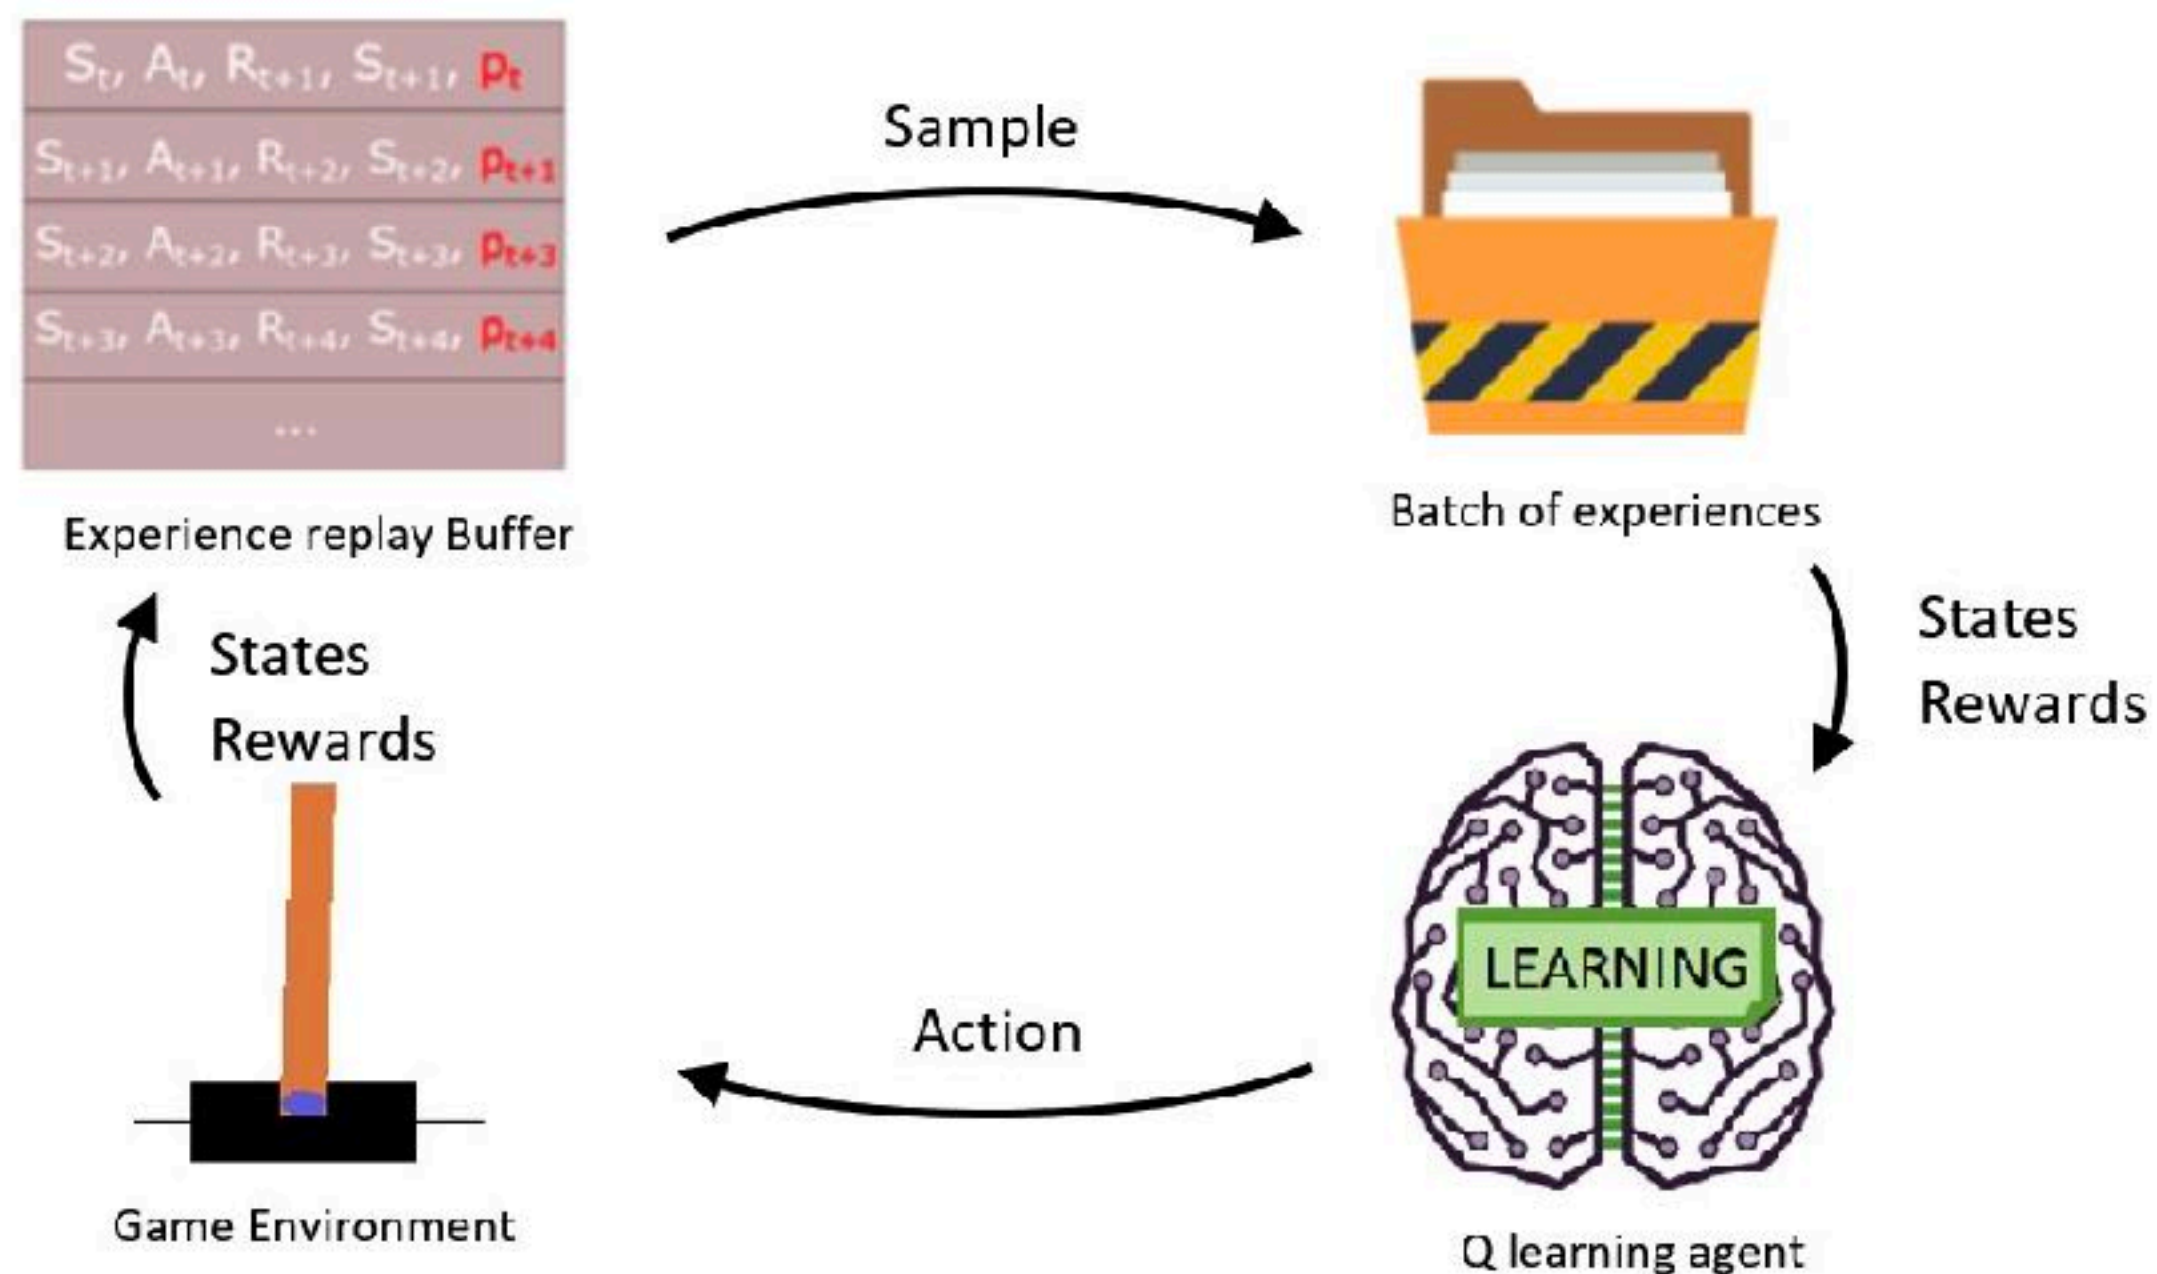
$$Q(s, a) = Q(s, a) + \alpha(R + \gamma \max_{a'} Q(s', a') - Q(s, a))$$

**Q-Network** - selects & evaluates actions



# DQN = Deep Q-Network

## Experience replay



```
def replay(self):
    if len(self.memory) < self.train_start:
        return
    # Randomly sample minibatch from the memory
    minibatch = random.sample(self.memory, min(len(self.memory), self.batch_size))

    state = np.zeros((self.batch_size, self.state_size))
    next_state = np.zeros((self.batch_size, self.state_size))
    action, reward, done = [], [], []

    # do this before prediction
    # for speedup, this could be done on the tensor level
    # but easier to understand using a loop
    for i in range(self.batch_size):
        state[i] = minibatch[i][0]
        action.append(minibatch[i][1])
        reward.append(minibatch[i][2])
        next_state[i] = minibatch[i][3]
        done.append(minibatch[i][4])

    # do batch prediction to save speed
    target = self.model.predict(state)
    target_next = self.model.predict(next_state)

    for i in range(self.batch_size):
        # correction on the Q value for the action used
        if done[i]:
            target[i][action[i]] = reward[i]
        else:
            # Standard - DQN
            # DQN chooses the max Q value among next actions
            # selection and evaluation of action is on the target Q Network
            # Q_max = max_a' Q_target(s', a')
            target[i][action[i]] = reward[i] + self.gamma * (np.amax(target_next[i]))

    # Train the Neural Network with batches
    self.model.fit(state, target, batch_size=self.batch_size, verbose=0)
```

# DQN = Deep Q-Network

## Training (one network) + Experience Replay

```
trainer = env.train([None, against])
observations = trainer.reset()
while not done:

    # take an action and store outcome
    action = TrainNet.get_action(observations, epsilon)
    prev_observations = observations
    observations, reward, done, _ = env.step(action)

    # Adding experience into buffer
    exp = {'s': prev_observations, 'a': action, 'r': reward,
          's2': observations, 'done': done}
    TrainNet.add_experience(exp)

    # Train the training model by using experiences in buffer
    # and the target model
    TrainNet.train(TargetNet)



    iter += 1
    if iter % copy_step == 0:
        # Update the weights of the target model after
        # reaching copy interval
        TargetNet.copy_weights(TrainNet)

return reward
```

# DQN = Deep Q-Network

Training (one network) + Experience Replay Buffer

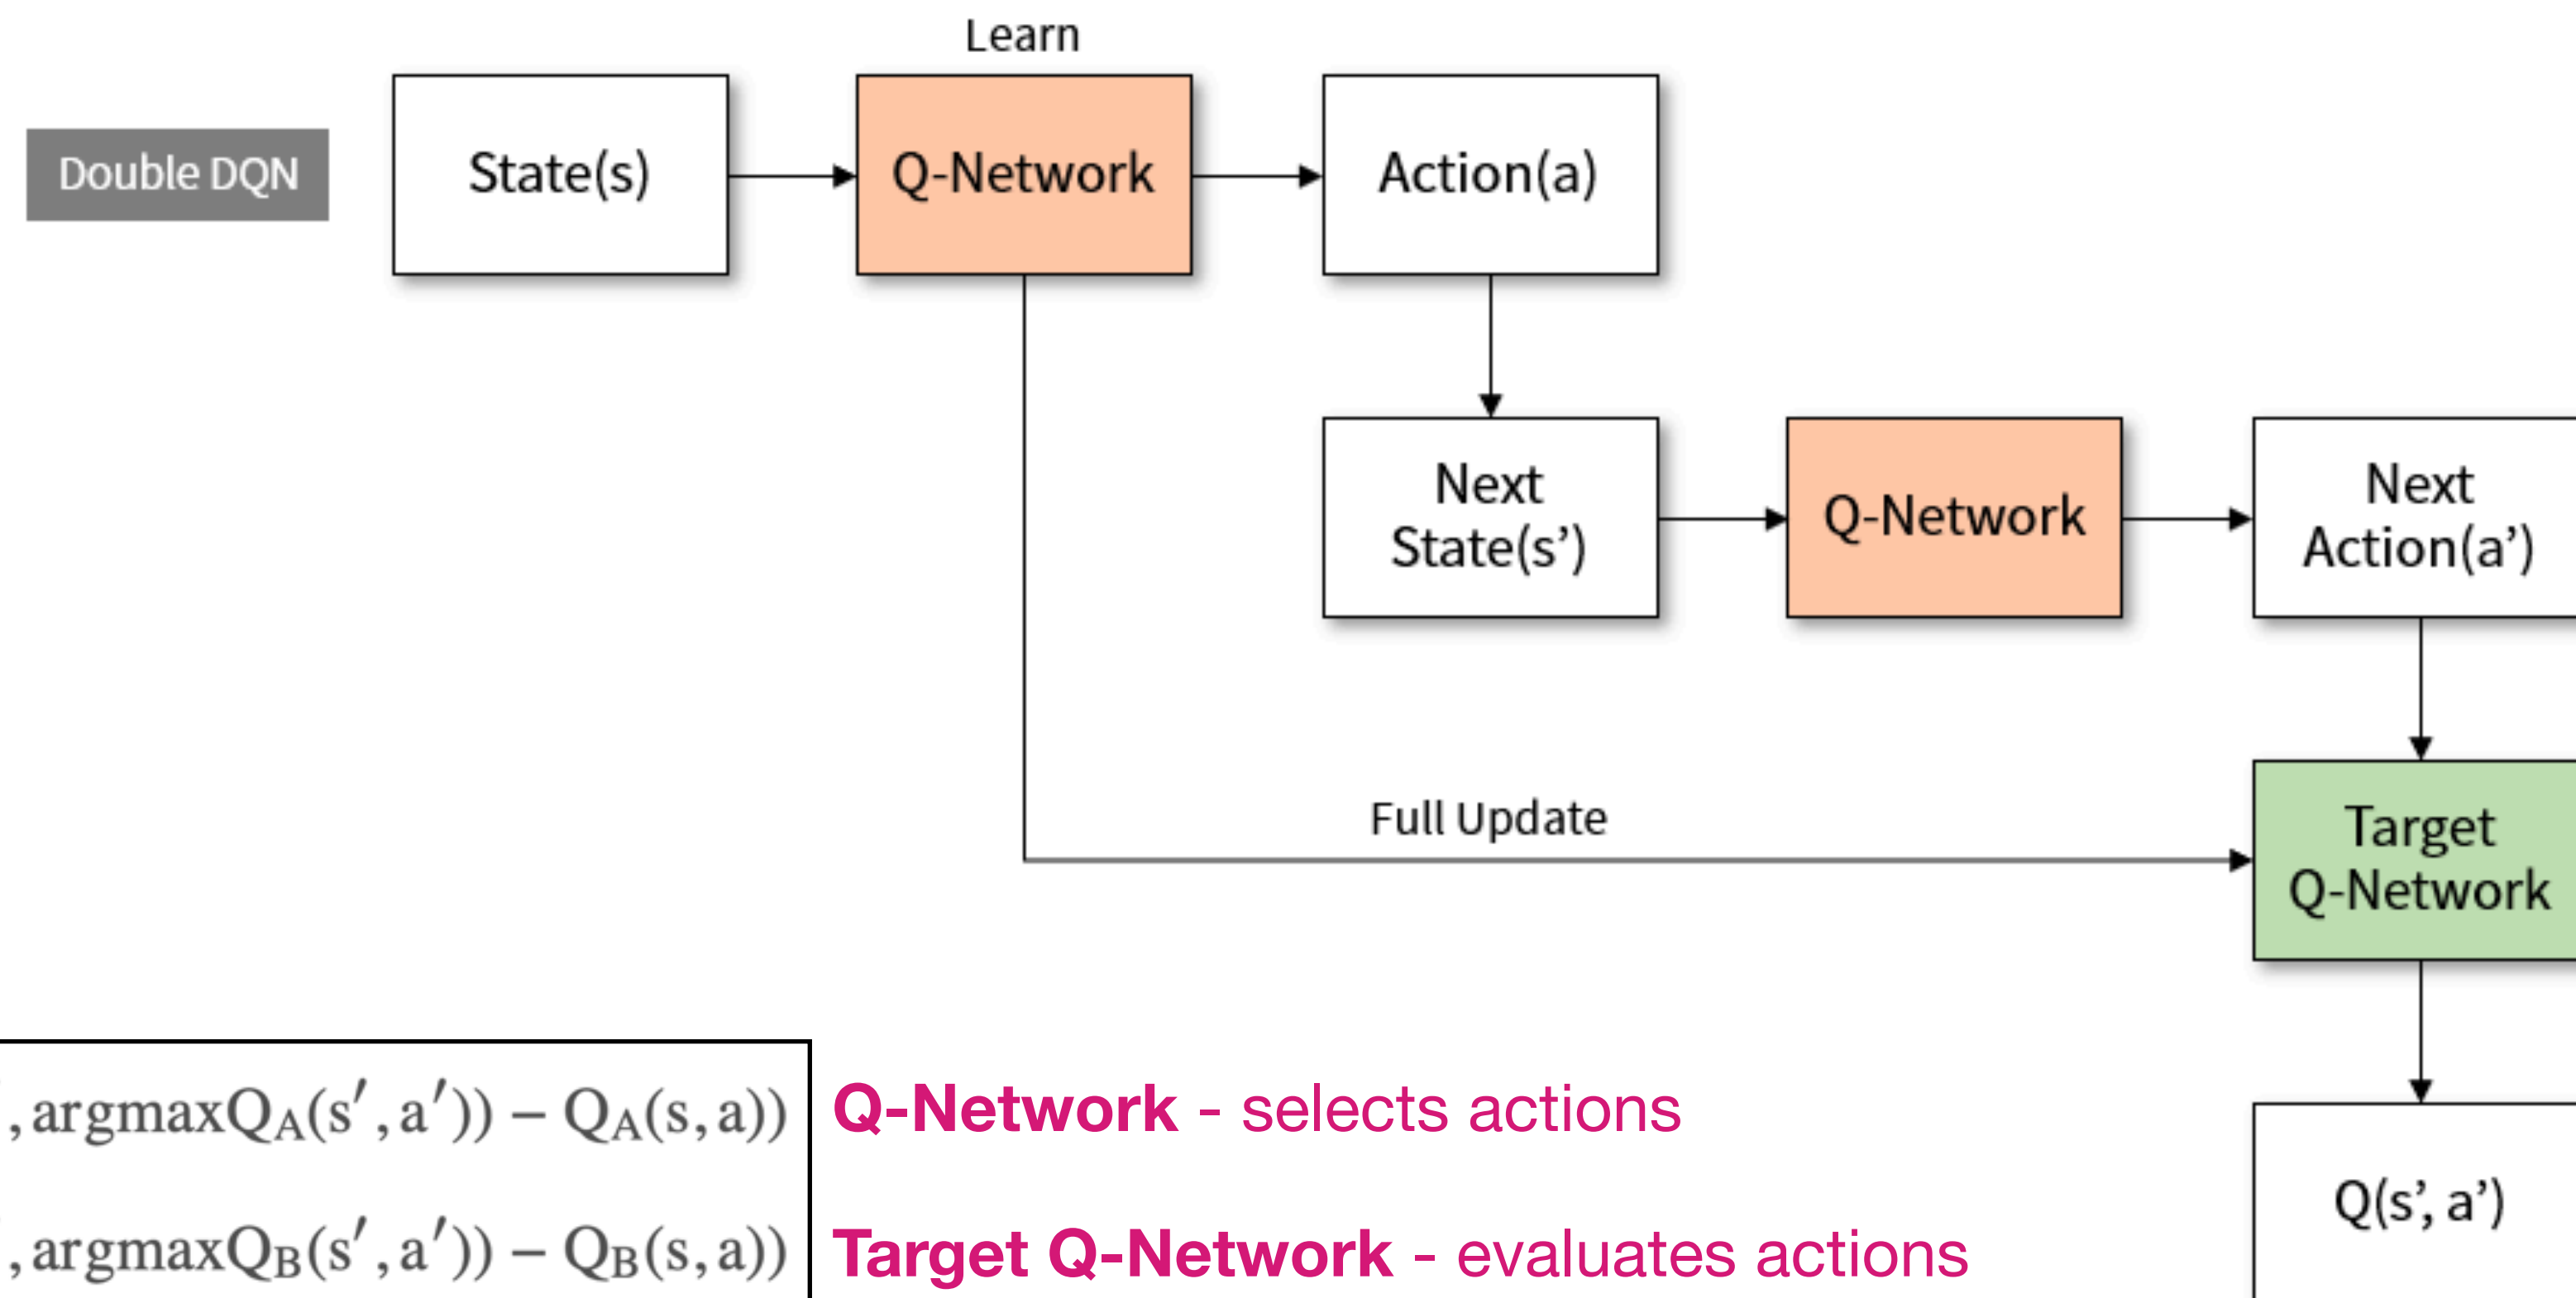
Average Reward





# Double DQN (2nd gen)

Two DQN networks converge faster, Target network's weights are updated once per N epochs



$$Q_A(s, a) = Q_A(s, a) + \alpha(R + \gamma Q_B(s', \operatorname{argmax} Q_A(s', a')) - Q_A(s, a))$$

$$Q_B(s, a) = Q_B(s, a) + \alpha(R + \gamma Q_A(s', \operatorname{argmax} Q_B(s', a')) - Q_B(s, a))$$

**Q-Network** - selects actions

**Target Q-Network** - evaluates actions

# Double DQN

Q-Network - selects actions

Target Q-Network - evaluates actions

$$Q_A(s, a) = Q_A(s, a) + \alpha(R + \gamma Q_B(s', \arg\max_{a'} Q_A(s', a')) - Q_A(s, a))$$

$$Q_B(s, a) = Q_B(s, a) + \alpha(R + \gamma Q_A(s', \arg\max_{a'} Q_B(s', a')) - Q_B(s, a))$$

```
def replay(self):
    if len(self.memory) < self.train_start:
        return
    # Randomly sample minibatch from the memory
    minibatch = random.sample(self.memory, min(self.batch_size, self.batch_size))

    state = np.zeros((self.batch_size, self.state_size))
    next_state = np.zeros((self.batch_size, self.state_size))
    action, reward, done = [], [], []

    # do this before prediction
    # for speedup, this could be done on the tensor level
    # but easier to understand using a loop
    for i in range(self.batch_size):
        state[i] = minibatch[i][0]
        action.append(minibatch[i][1])
        reward.append(minibatch[i][2])
        next_state[i] = minibatch[i][3]
        done.append(minibatch[i][4])

    # do batch prediction to save speed
    target = self.model.predict(state)
    target_next = self.model.predict(next_state)
    target_val = self.target_model.predict(next_state)

    for i in range(len(minibatch)):
        # correction on the Q value for the action used
        if done[i]:
            target[i][action[i]] = reward[i]
        else:
            if self.ddqn: # Double - DQN
                # current Q Network selects the action
                # a'_max = argmax_a' Q(s', a')
                a = np.argmax(target_next[i])
                # target Q Network evaluates the action
                # Q_max = Q_target(s', a'_max)
                target[i][action[i]] = reward[i] + self.gamma * (target_val[i][a])
            else: # Standard - DQN
                # DQN chooses the max Q value among next actions
                # selection and evaluation of action is on the target Q Network
                # Q_max = max_a' Q_target(s', a')
                target[i][action[i]] = reward[i] + self.gamma * (np.amax(target_next[i]))

    # Train the Neural Network with batches
    self.model.fit(state, target, batch_size=self.batch_size, verbose=0)
```



# Double DQN

**Q-Network - selects actions**

**Target Q-Network - evaluates actions**

- The Bellman equation used to calculate the Q values to update the online network follows the equation:

```
value = reward + discount_factor *  
target_network.predict(next_state)  
[argmax(online_network.predict(next_state))]
```

- The Bellman equation used to calculate the Q value updates in the original DQN is:

```
value = reward + discount_factor *  
max(target_network.predict(next_state))
```

# Double DQN - Soft Target Update

$\text{target\_weights} = \text{target\_weights} * (1 - \text{TAU}) + \text{q\_weights} * \text{TAU}$  where  $0 < \text{TAU} < 1$

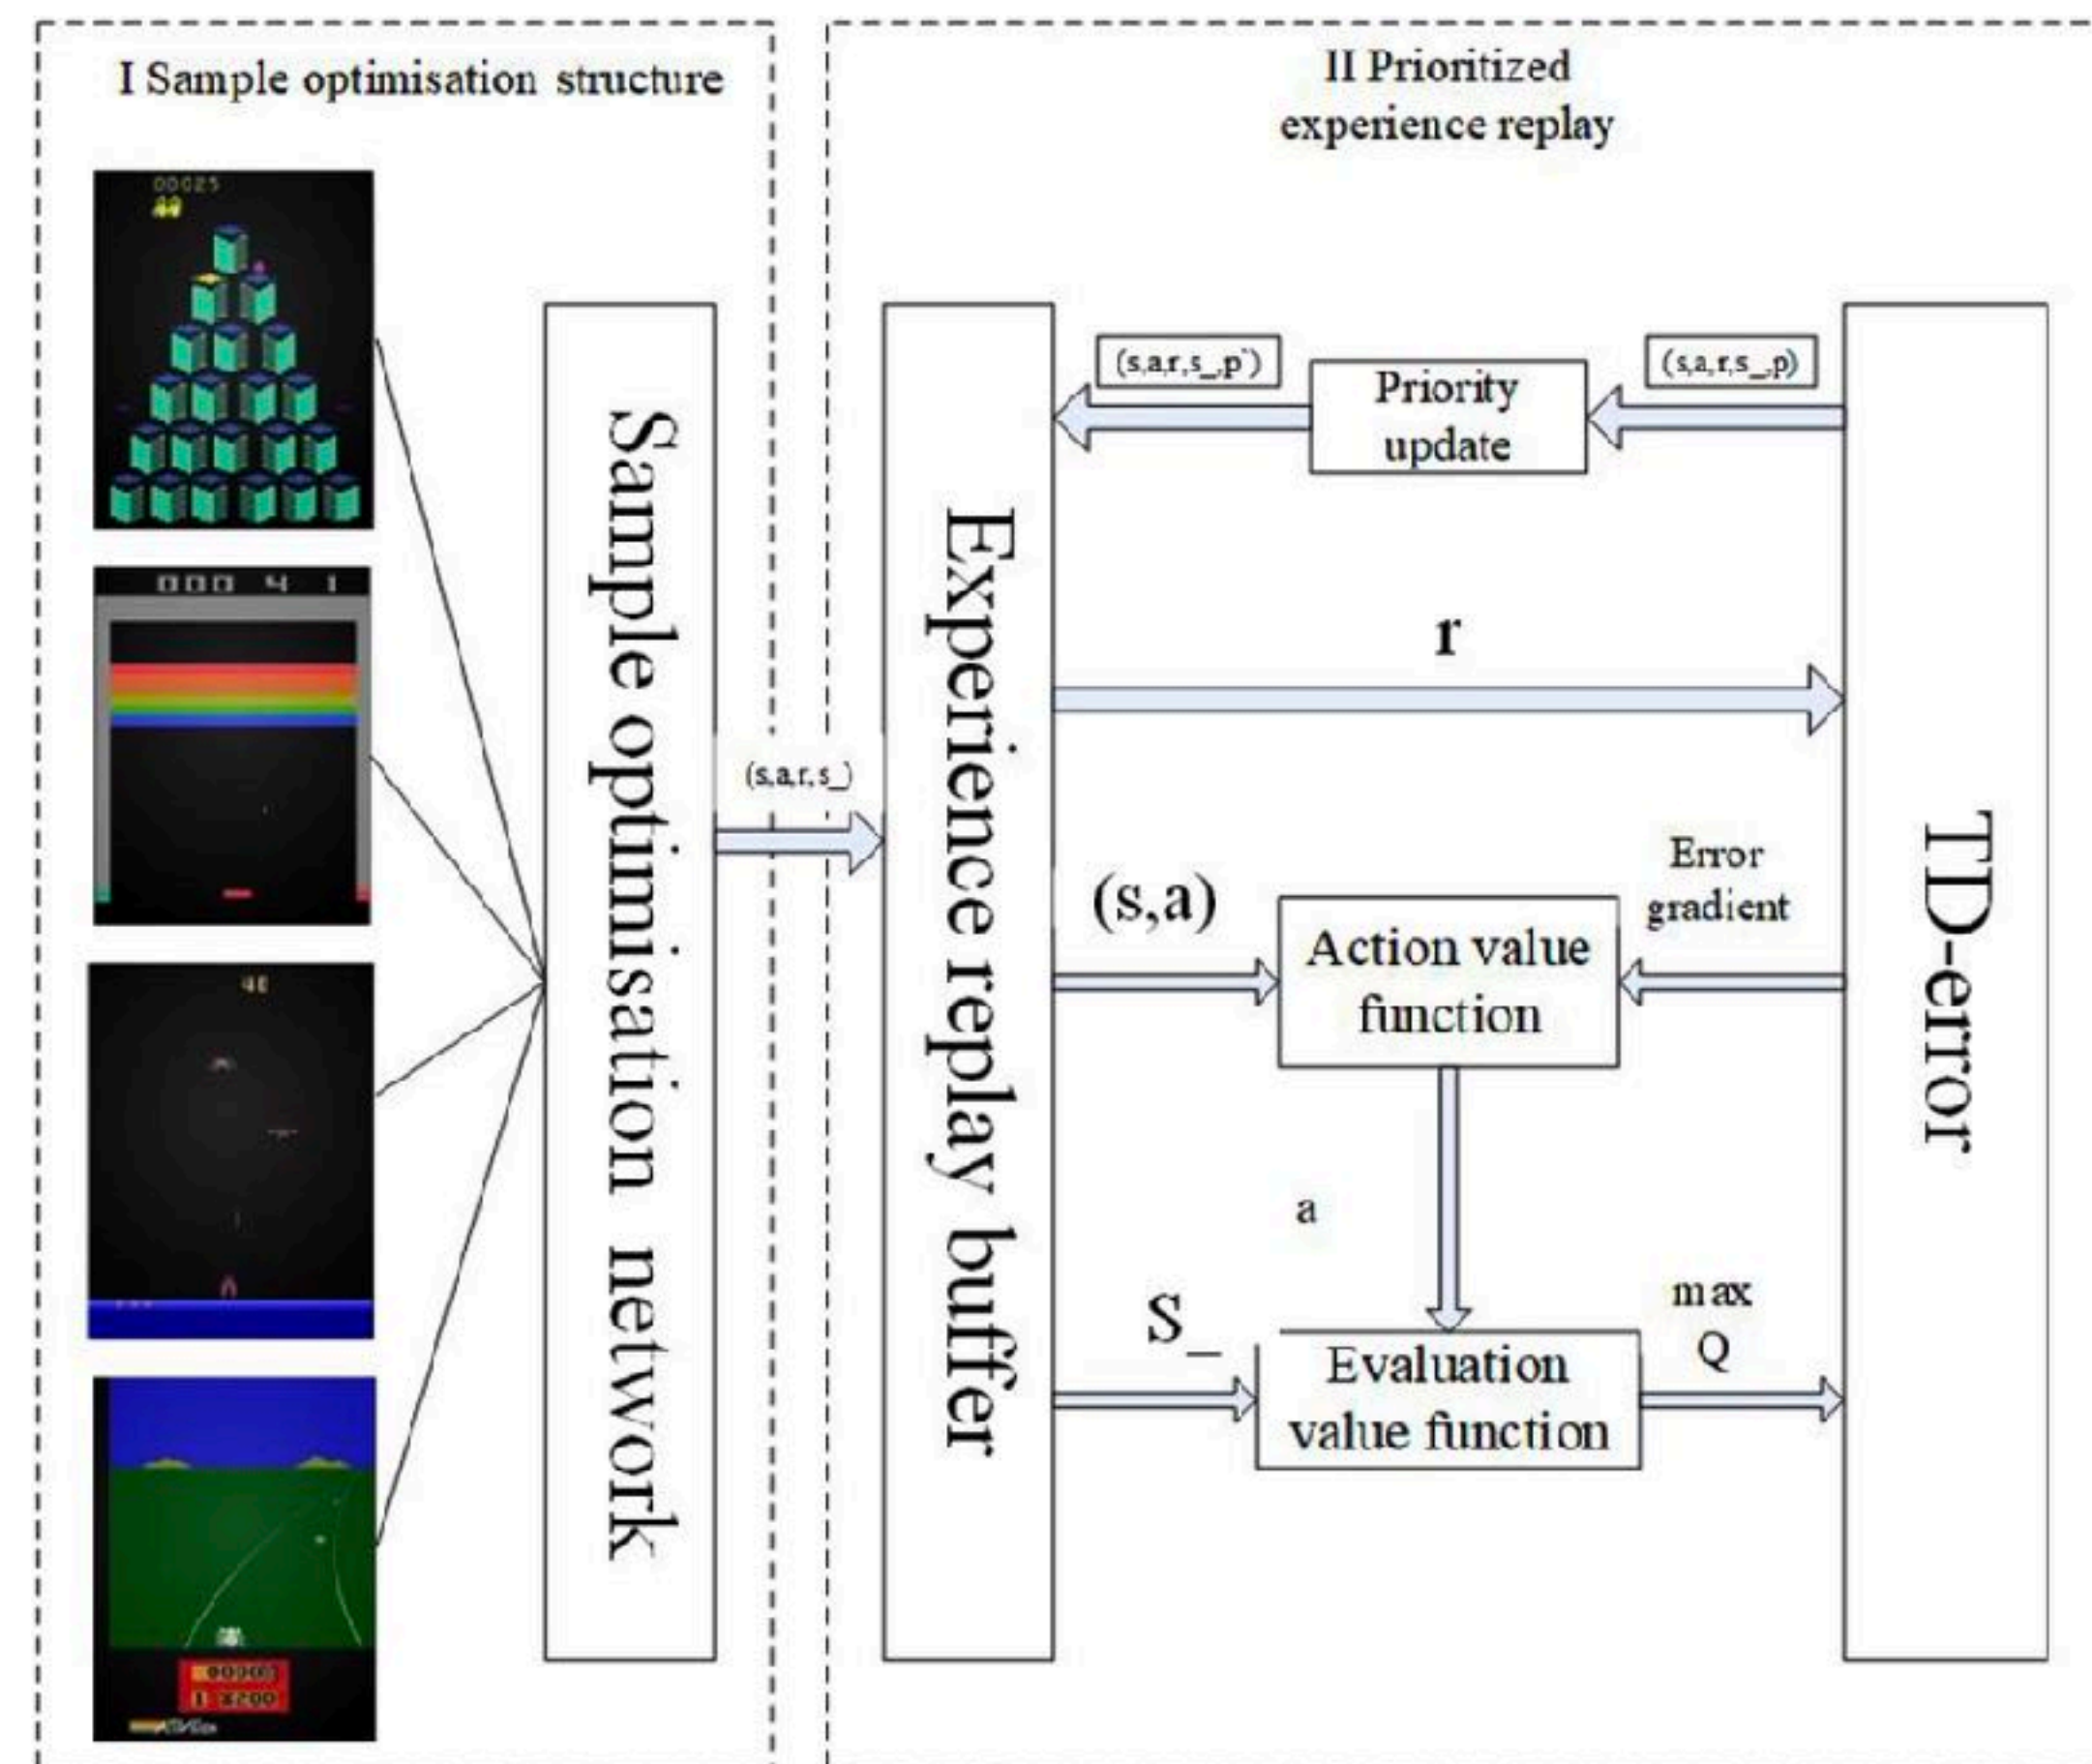
```
def update_target_model(self):
    if not self.Soft_Update and self.ddqn:
        self.target_model.set_weights(self.model.get_weights())
        return
    if self.Soft_Update and self.ddqn:
        q_model_theta = self.model.get_weights()
        target_model_theta = self.target_model.get_weights()
        counter = 0
        for q_weight, target_weight in zip(q_model_theta, target_model_theta):
            target_weight = target_weight * (1 - self.TAU) + q_weight * self.TAU
            target_model_theta[counter] = target_weight
            counter += 1
        self.target_model.set_weights(target_model_theta)
```

# Prioritised Replay (2nd gen)

For TD-learning.

Order of replying updates could help speed up learning. Priority of a tuple  $s_i - a_i - r_i - s_{i+1}$  is proportional to TD error

More informative experience replay -> faster training



# Duelling DQN (2nd gen)

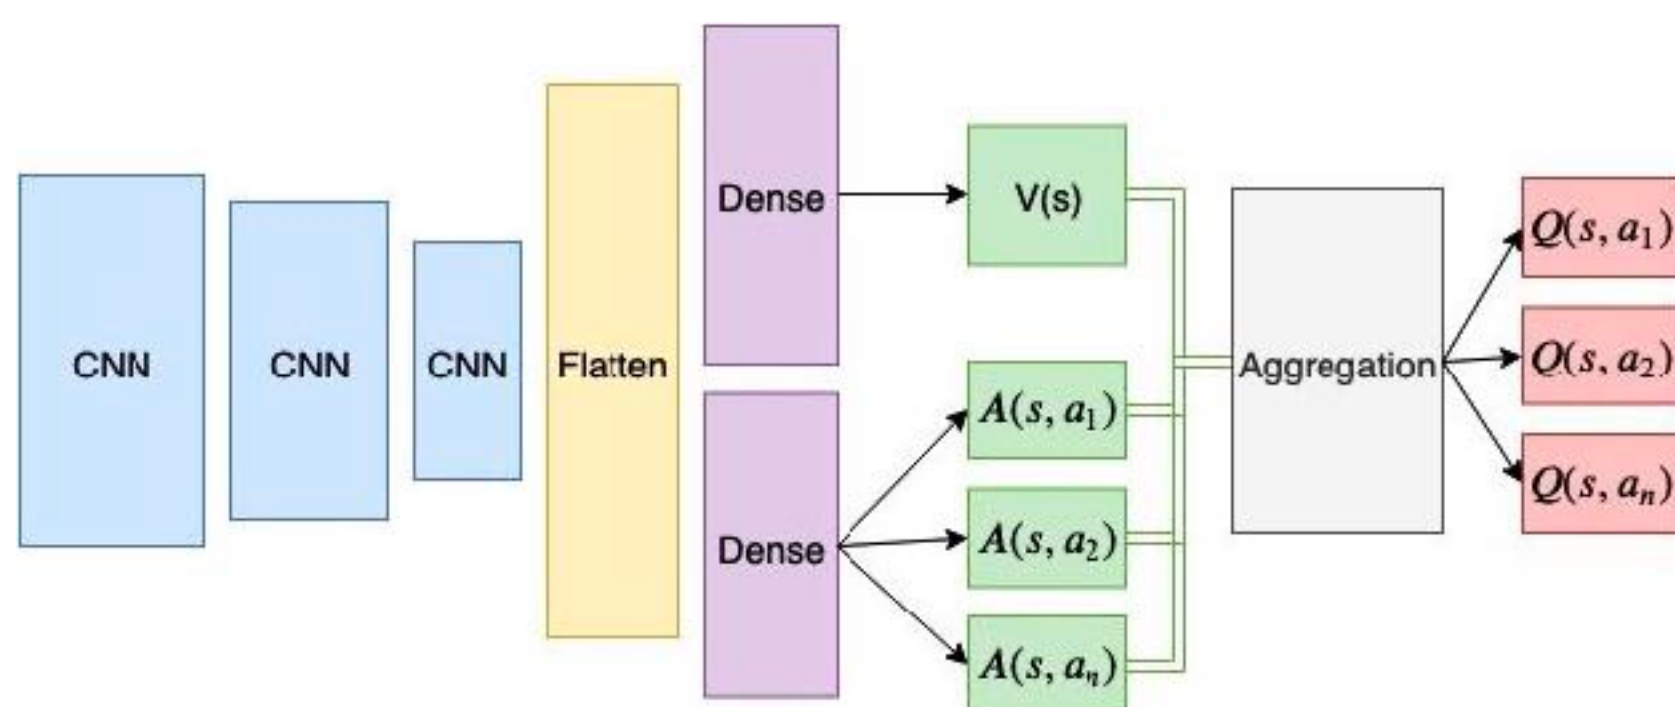
## Why

- TDQN Networks tend to overestimate rewards in noisy environments, leading to non-optimal training outcomes;
- The moving target problem is that the same network is responsible for choosing and evaluating actions, leading to training instability.
- **Double Dueling DQN**: the evaluation of the Q function implicitly calculates two quantities:
  - $V(s)$**  – the value of being in state  $s$ ;
  - $A(s, a)$**  – the advantage of taking action  $a$  in state  $s$ .



# Duelling DQN (2nd gen)

Has two heads - for  $V(s)$  and  $A(s,a)$  estimation



```
def OurModel(input_shape, action_space, dueling):
    X_input = Input(input_shape)
    X = X_input

    # 'Dense' is the basic form of a neural network layer
    # Input Layer of state size(4) and Hidden Layer with 512 nodes
    X = Dense(512, input_shape=input_shape, activation="relu", kernel_initializer='he_uniform')(X)

    # Hidden layer with 256 nodes
    X = Dense(256, activation="relu", kernel_initializer='he_uniform')(X)

    # Hidden layer with 64 nodes
    X = Dense(64, activation="relu", kernel_initializer='he_uniform')(X)

    if dueling:
        state_value = Dense(1, kernel_initializer='he_uniform')(X)
        state_value = Lambda(lambda s: K.expand_dims(s[:, 0], -1), output_shape=(action_space,))(state_value)

        action_advantage = Dense(action_space, kernel_initializer='he_uniform')(X)
        action_advantage = Lambda(lambda a: a[:, :] - K.mean(a[:, :], keepdims=True), output_shape=(action_space,))(action_advantage)

        X = Add()([state_value, action_advantage])
    else:
        # Output Layer with # of actions: 2 nodes (left, right)
        X = Dense(action_space, activation="linear", kernel_initializer='he_uniform')(X)

    model = Model(inputs = X_input, outputs = X, name='CartPole Dueling DDQN model')
    model.compile(loss="mean_squared_error", optimizer=RMSprop(lr=0.00025, rho=0.95, epsilon=0.01), metrics=["accuracy"])

    model.summary()
    return model
```



**Ape-X** =  
*Distributed* Prioritized Experience  
Replay  
+  
DQN / Double DQN / Duelling  
DQN

# Ape-X config for Ray RLlib

Example: [https://github.com/ray-project/ray/blob/master/rllib/algorithms/apex\\_dqn/apex\\_dqn.py](https://github.com/ray-project/ray/blob/master/rllib/algorithms/apex_dqn/apex_dqn.py)

```
def __init__(self, algo_class=None):
    """Initializes a ApexConfig instance."""
    super().__init__(algo_class=algo_class or ApexDQN)

    # fmt: off
    # __sphinx_doc_begin__
    # APEX-DQN settings overriding DQN ones:
    # .training()
    self.optimizer = merge_dicts(
        DQNConfig().optimizer, {
            "max_weight_sync_delay": 400,
            "num_replay_buffer_shards": 4,
            "debug": False
        })
    self.n_step = 3
    self.train_batch_size = 512
    self.target_network_update_freq = 500000
    self.training_intensity = 1
    # Number of timesteps to collect from rollout workers before we start
    # sampling from replay buffers for learning. Whether we count this in agent
    # steps or environment steps depends on config["multiagent"]["count_steps_by"].
    self.num_steps_sampled_before_learning_starts = 50000

    self.max_requests_in_flight_per_replay_worker = float("inf")
    self.timeout_s_sampler_manager = 0.0
    self.timeout_s_replay_manager = 0.0
    # APEX-DQN is using a distributed (non local) replay buffer.
    self.replay_buffer_config = {
        "no_local_replay_buffer": True,
        # Specify prioritized replay by supplying a buffer type that supports
        # prioritization
        "type": "MultiAgentPrioritizedReplayBuffer",
        "capacity": 2000000,
```

**Rainbow?**

**= a successful combinations of  
DeepRL improvements**

# Rainbow

Example: [https://github.com/ray-project/ray/blob/master/rllib/tuned\\_examples/dqn/pong-rainbow.yaml](https://github.com/ray-project/ray/blob/master/rllib/tuned_examples/dqn/pong-rainbow.yaml)

```
1  pong-deterministic-rainbow:
2      env: PongDeterministic-v4
3      run: DQN
4      stop:
5          episode_reward_mean: 20
6      config:
7          num_atoms: 51
8          noisy: True
9          gamma: 0.99
10         lr: .0001
11         hiddens: [512]
12         rollout_fragment_length: 4
13         train_batch_size: 32
14         exploration_config:
15             epsilon_timesteps: 2
16             final_epsilon: 0.0
17         target_network_update_freq: 500
18         replay_buffer_config:
19             type: MultiAgentPrioritizedReplayBuffer
20             prioritized_replay_alpha: 0.5
21             capacity: 50000
22         num_steps_sampled_before_learning_starts: 10000
23         n_step: 3
24         gpu: True
25         model:
26             grayscale: True
27             zero_mean: False
28             dim: 42
29         # we should set compress_observations to True because few machines
30         # would be able to contain the replay buffers in memory otherwise
31         compress_observations: True
```

# Homework

## Assault + Gym Env + Ray RLlib

- Train & tune DQN
- Train & tune DQN + Experience Replay
- Train & tune Double DQN + Experience Replay
- Train & tune Double DQN + Prioritised Experience Replay
- Train & tune Double Duelling DQN + Prioritised Experience Replay
- Train & tune Ape-X: Double Duelling DQN + Distributed Prioritised Experience Replay
- Cheatsheet: [https://github.com/ray-project/ray/tree/master/rllib/tuned\\_examples](https://github.com/ray-project/ray/tree/master/rllib/tuned_examples)