

Playing Atari with Deep Reinforcement Learning

Part 2:

Policy Gradients, A2C, A3C, PPO

By Gurbych Oleksandr

Playing Atari with Deep Reinforcement Learning

Part 2:

Policy Gradients, A2C, A3C, PPO

By Gurbych Oleksandr

Intro

Policy Gradients

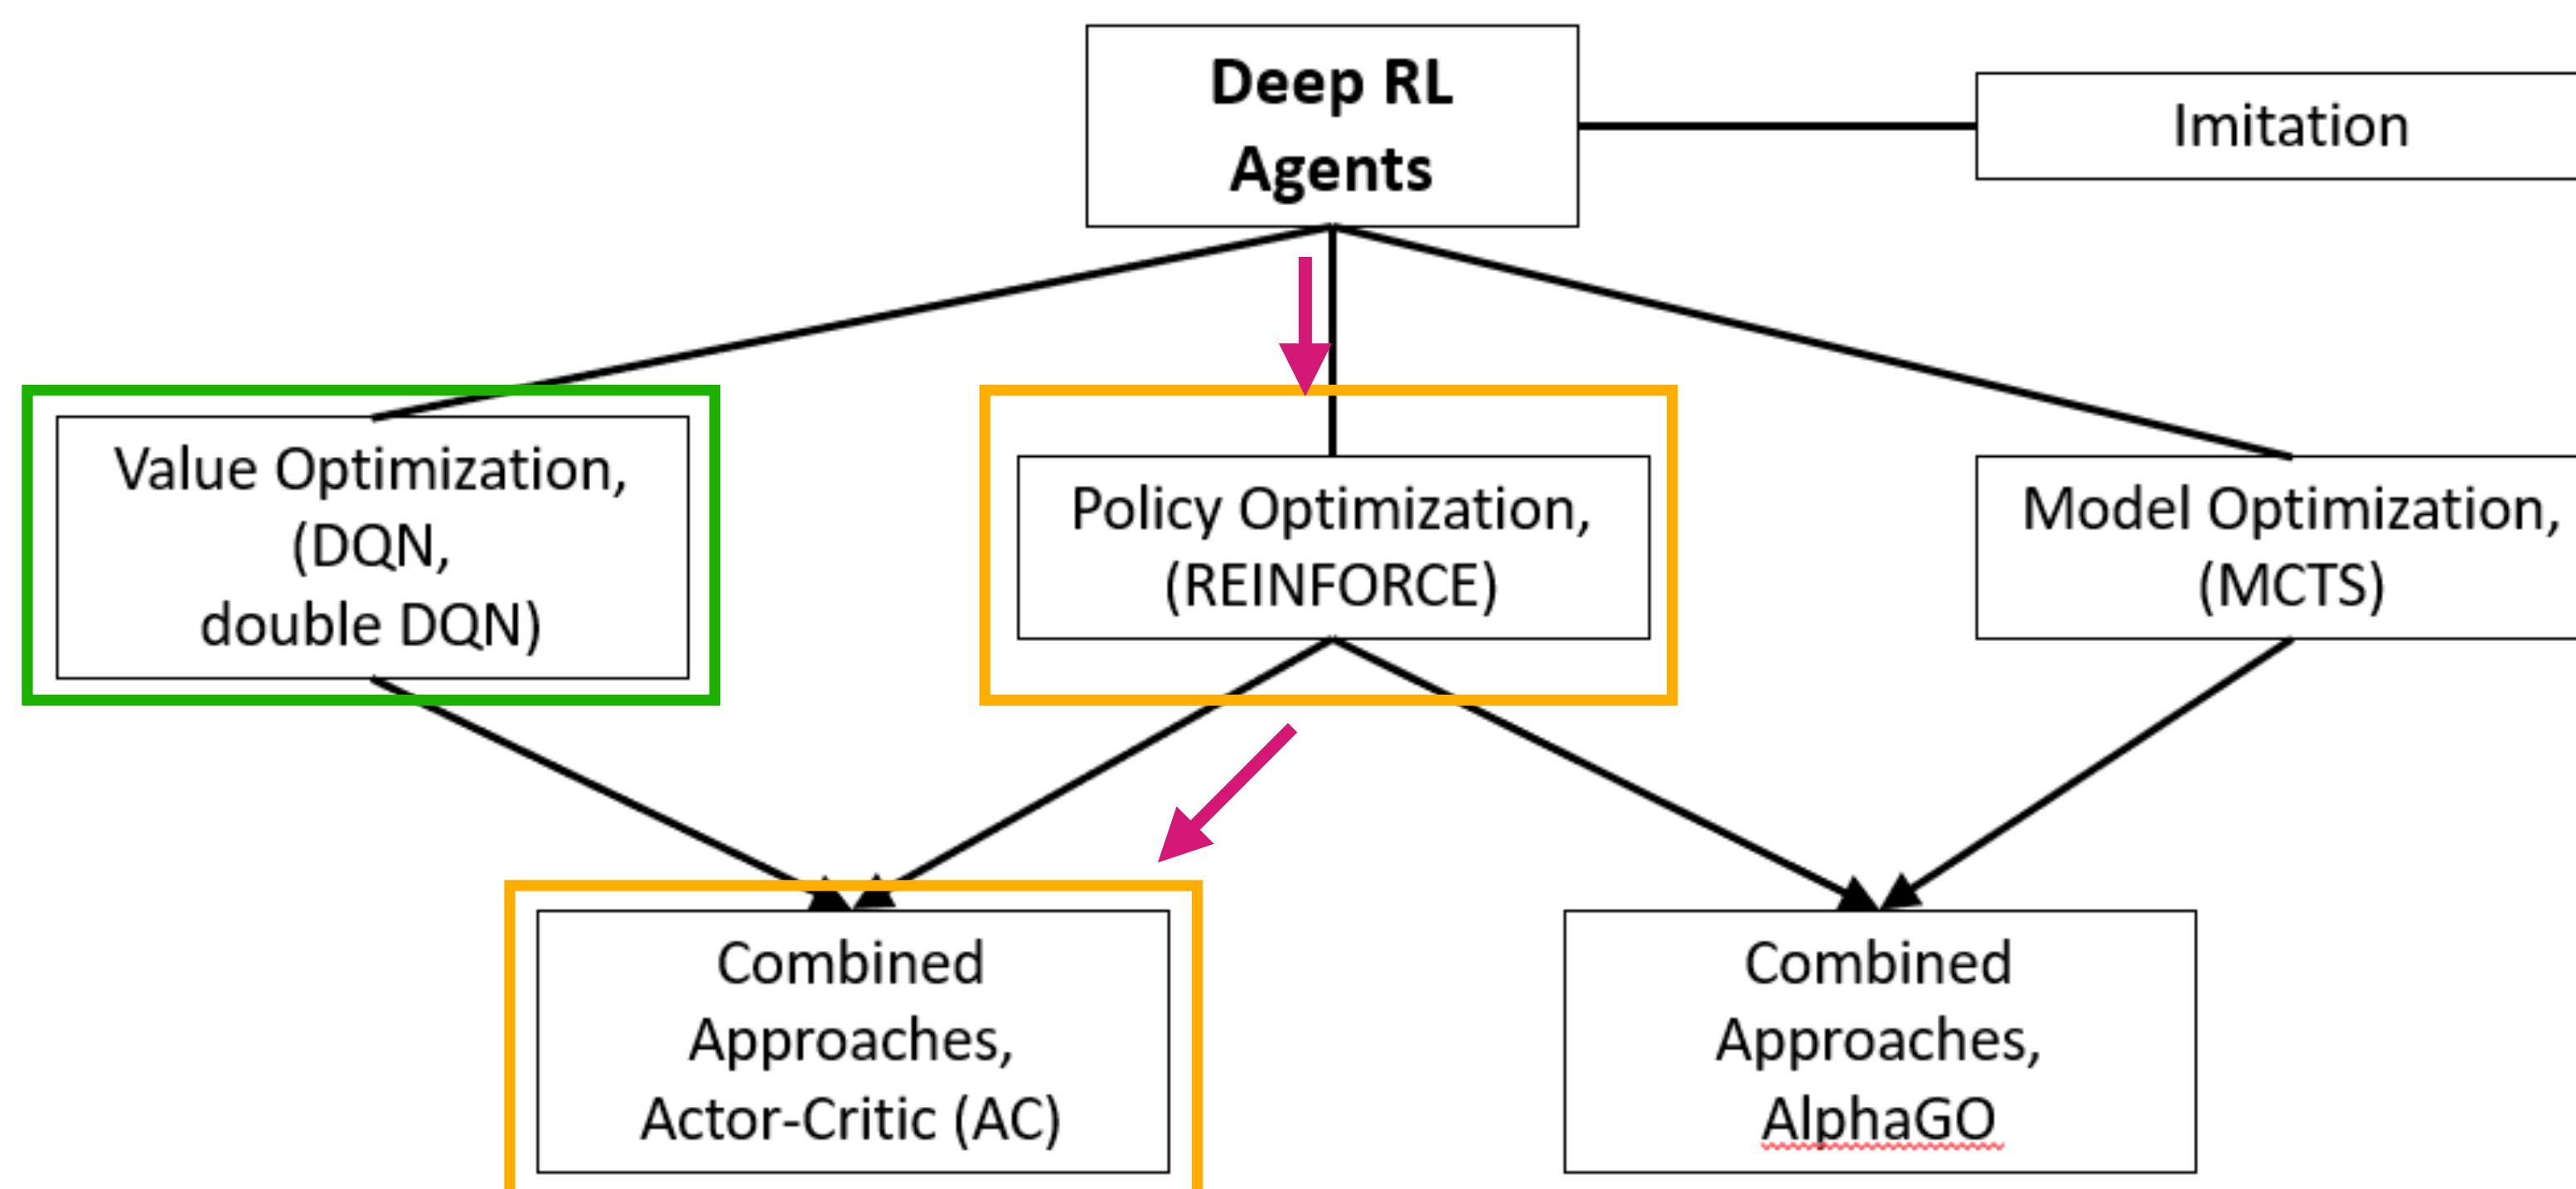
DQN Drawbacks

DQN's are comparatively simple and efficient but...

- Significant **oscillations while training**. This is because the choice of action may change dramatically for an arbitrarily small change in the estimated action values.
- Suppose the possible number of state-action pairs is relatively large in a given environment. In that case, **the Q-function can become highly complicated**, so it becomes intractable to estimate the optimal Q-value.
- Even in situations where finding Q is computationally tractable, **DQN's are not great at exploring** relative to some other approaches, so a DQN may not work correctly.

Other Types of Deep RL Agents

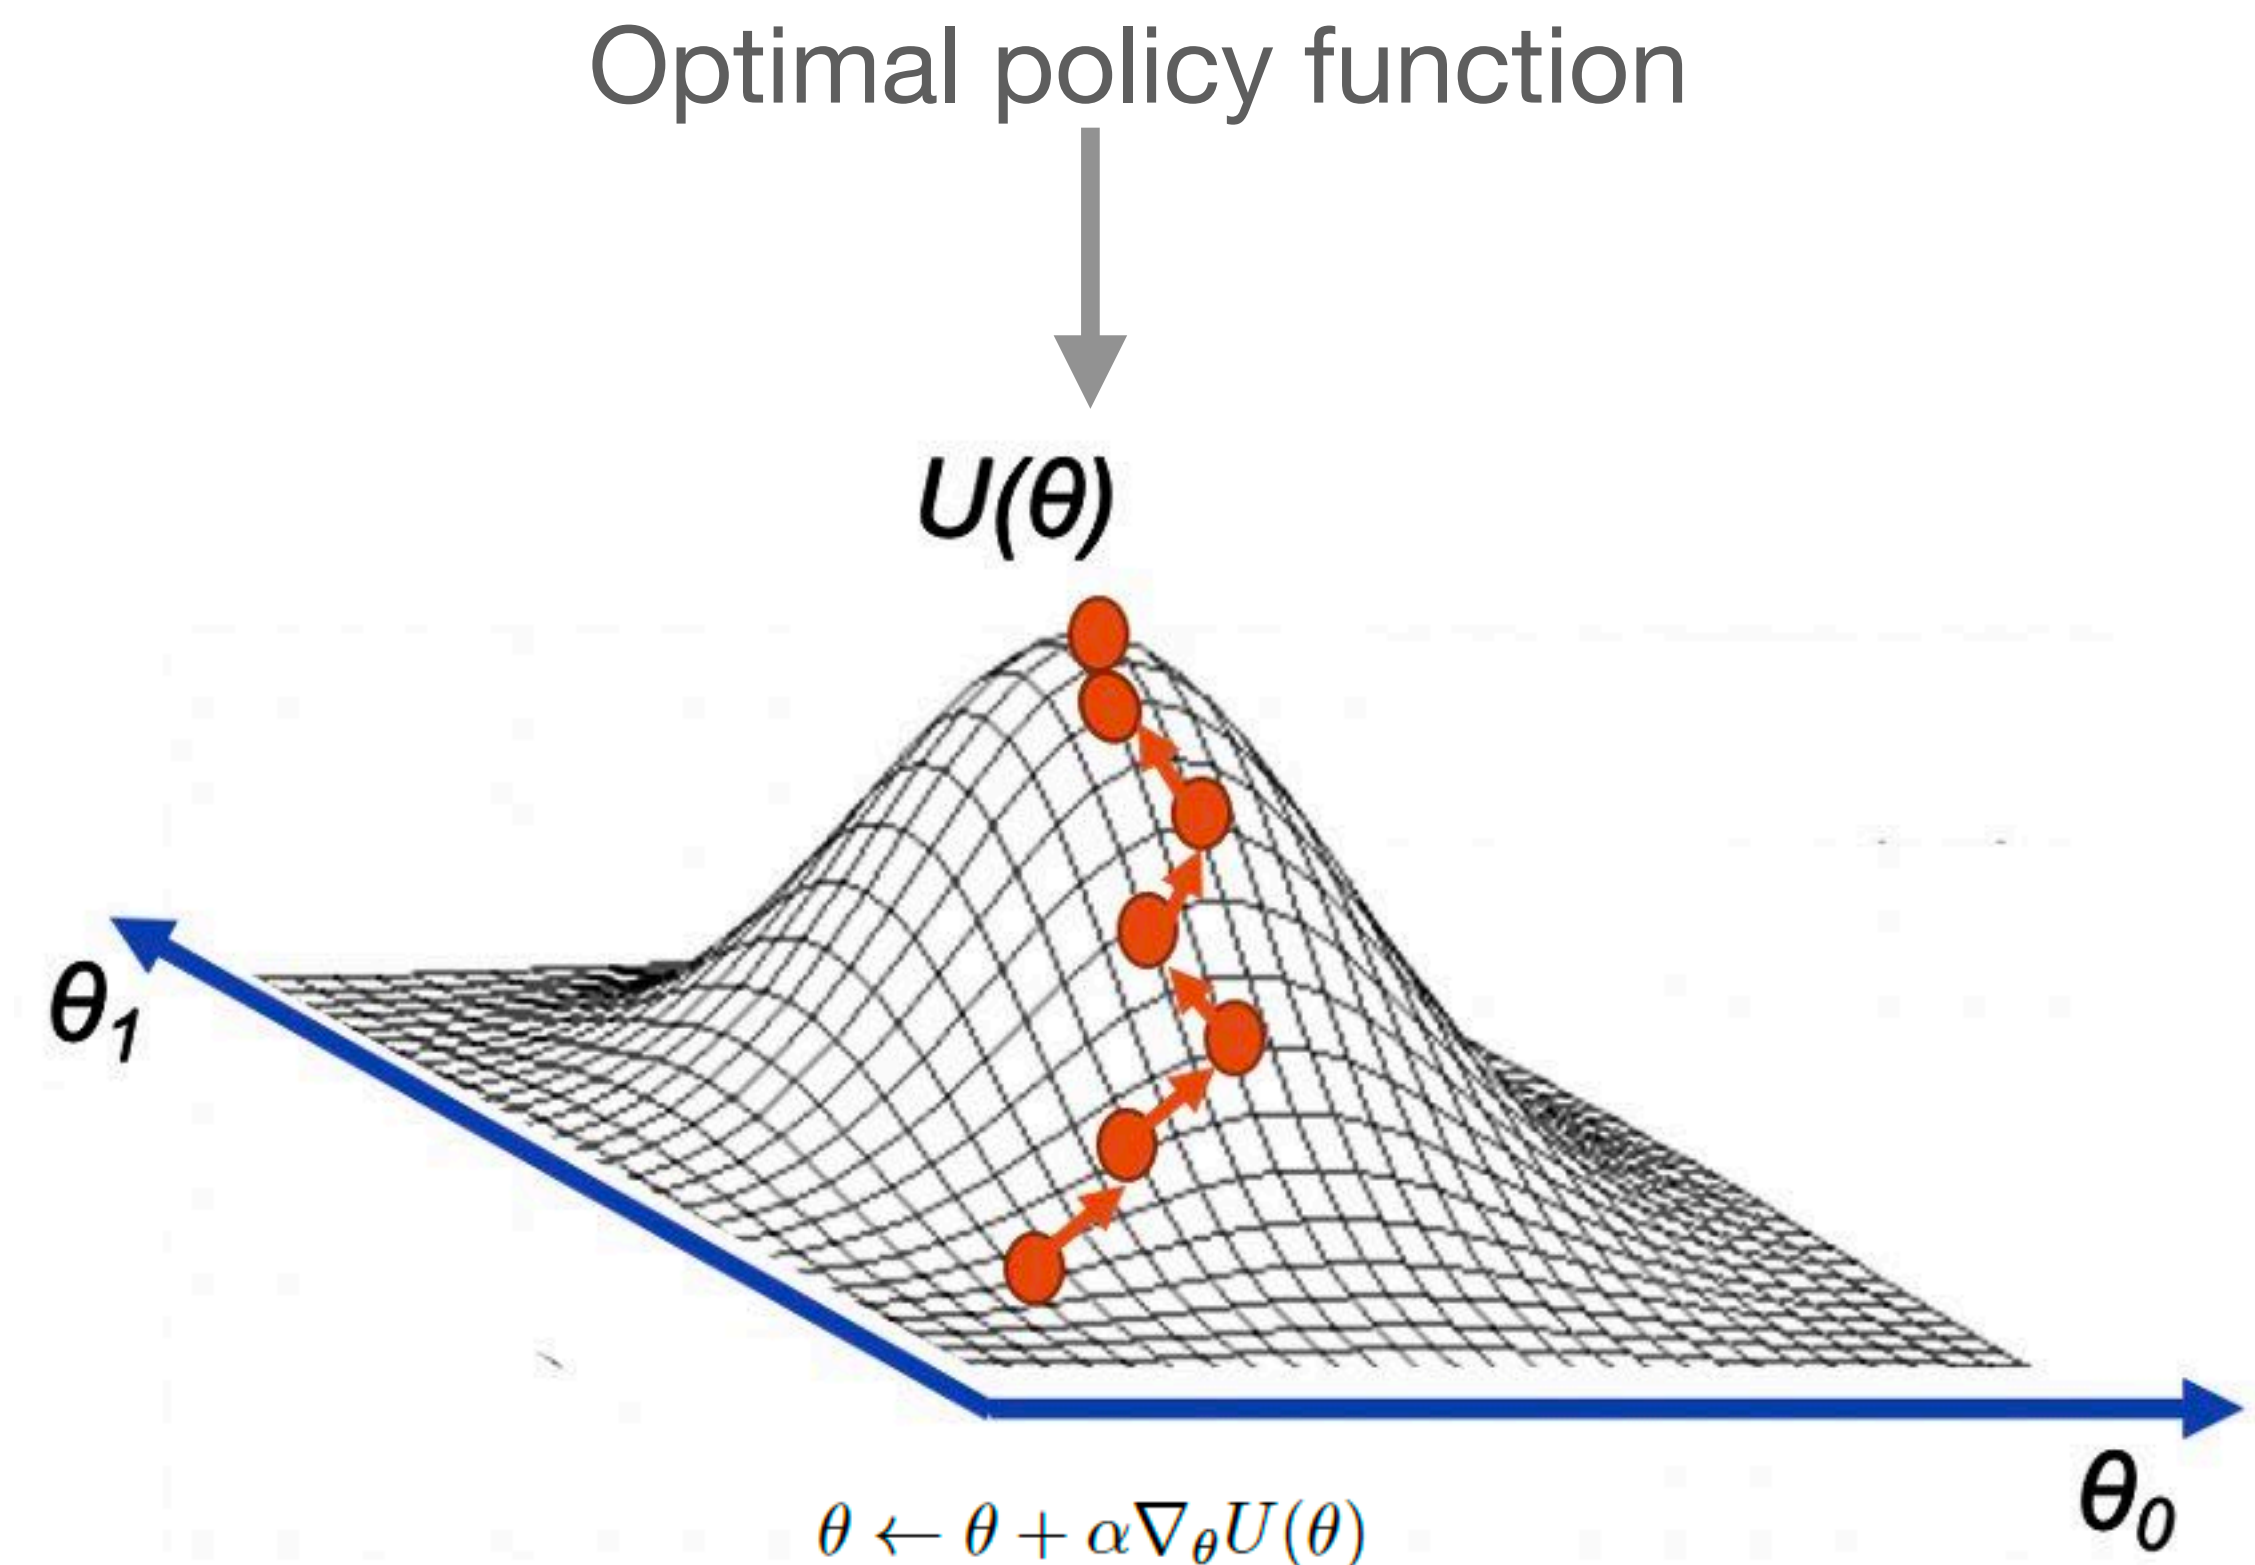
DQN's are comparatively simple and efficient but...



Value Optimisation vs Policy Optimisation

...implements the classic “agent-environment loop”


- A policy function π maps the state space \mathbf{S} to the action space \mathbf{A}
- DQN’s “policy” π is learned indirectly - by picking the (s,a) pairs with the largest Q-value - choosing “the best action”
- Policy gradient (PG) algorithms perform gradient ascent on π directly and learn the policy π directly



Value Optimisation vs Policy Optimisation

...implements the classic “agent-environment loop”

- PG algorithms **converge faster** and robuster than value optimisation algorithms such as DQN
- PGs are more **effective in large action spaces** or using continuous actions.
- Deep Q-learning assigns a score (maximum expected future reward) **for every possible action**, within each time step, given the current state. But what if we have endless possibilities for action?



Policy Gradient

...fit states to actions directly

Input: a differentiable policy parameterization $\pi(a|s, \theta)$

Algorithm parameter: step size $\alpha > 0$

Initialize the policy parameter θ at random

(1) Use the policy π_θ to collect a trajectory $\tau = (s_0, a_0, r_1, s_1, a_1, r_2, s_2, \dots, a_H, r_{H+1}, s_{H+1})$

(2) Estimate the Return for trajectory τ : $R(\tau) = (G_0, G_1, \dots, G_H)$
where G_k is the expected return for transition k :

$$G_k \leftarrow \sum_{t=k+1}^{H+1} \gamma^{t-k-1} R_t$$

(3) Use the trajectory τ to estimate the gradient $\nabla_\theta U(\theta)$

$$\nabla_\theta U(\theta) \leftarrow \sum_{t=0}^H \nabla_\theta \log \pi_\theta(a_t|s_t) G_t$$

(4) Update the weights θ of the policy

$$\theta \leftarrow \theta + \alpha \nabla_\theta U(\theta)$$

(5) Loop over steps 1-5 until not converged

$$\Delta J(Q) = E_\tau \left[\sum_{t=0}^{T-1} \nabla_Q \log \pi_Q(a_t, s_t) G_t \right]$$

Policy function
(differentiable)

Policy Gradient

Drawbacks

$$\Delta J(Q) = E_T \left[\sum_{t=0}^{T-1} \nabla_Q \log \pi_Q(a_t, s_t) G_t \right] \longleftarrow \text{Policy function}$$

- The PG algorithm updates the policy using Monte Carlo (i.e., taking random samples) =>
each training trajectory can be very different (1) =>
high variability in log probs and cumulative reward =>
noisy gradients => **unstable learning** => non-optimal policy distribution learned
- (1) => **trajectories with with cumulative reward 0** => PG algorithm doesn't improve there


Policy Gradient

Drawbacks

$$\Delta J(Q) = E_T \left[\sum_{t=0}^{T-1} \nabla_Q \log \pi_Q(a_t, s_t) G_t \right]$$

← Policy function

- **Unstable**
- **Slow convergence**




Section 2

Advantage Actor-Critic (A2C)

Advantage Actor-Critic (A2C)

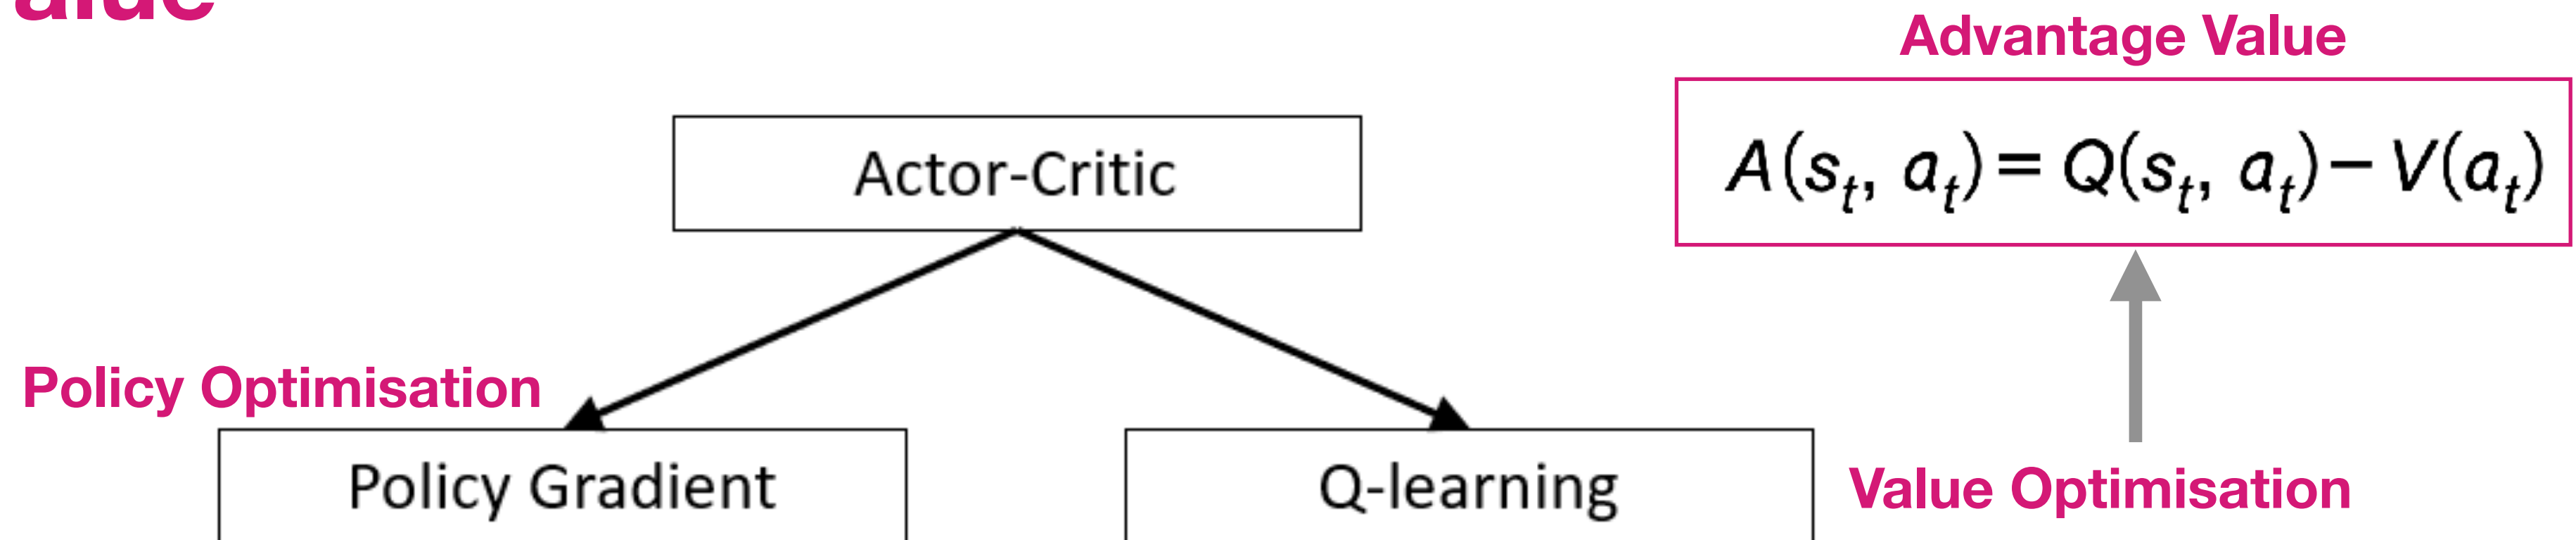
...implements the classic “agent-environment loop”



- **Actor**: a PG algorithm that decides on an action to take;
- **Critic**: Q-learning algorithm that critiques the action that the Actor selected, providing feedback on how to adjust. It can take advantage of efficiency tricks in Q-learning, such as memory replay.
- **A2C can solve a broader range of problems** than DQN
- **A2C has a lower variance in performance** relative to a pure PG
- Sampling inefficient

Advantage Actor-Critic (A2C)

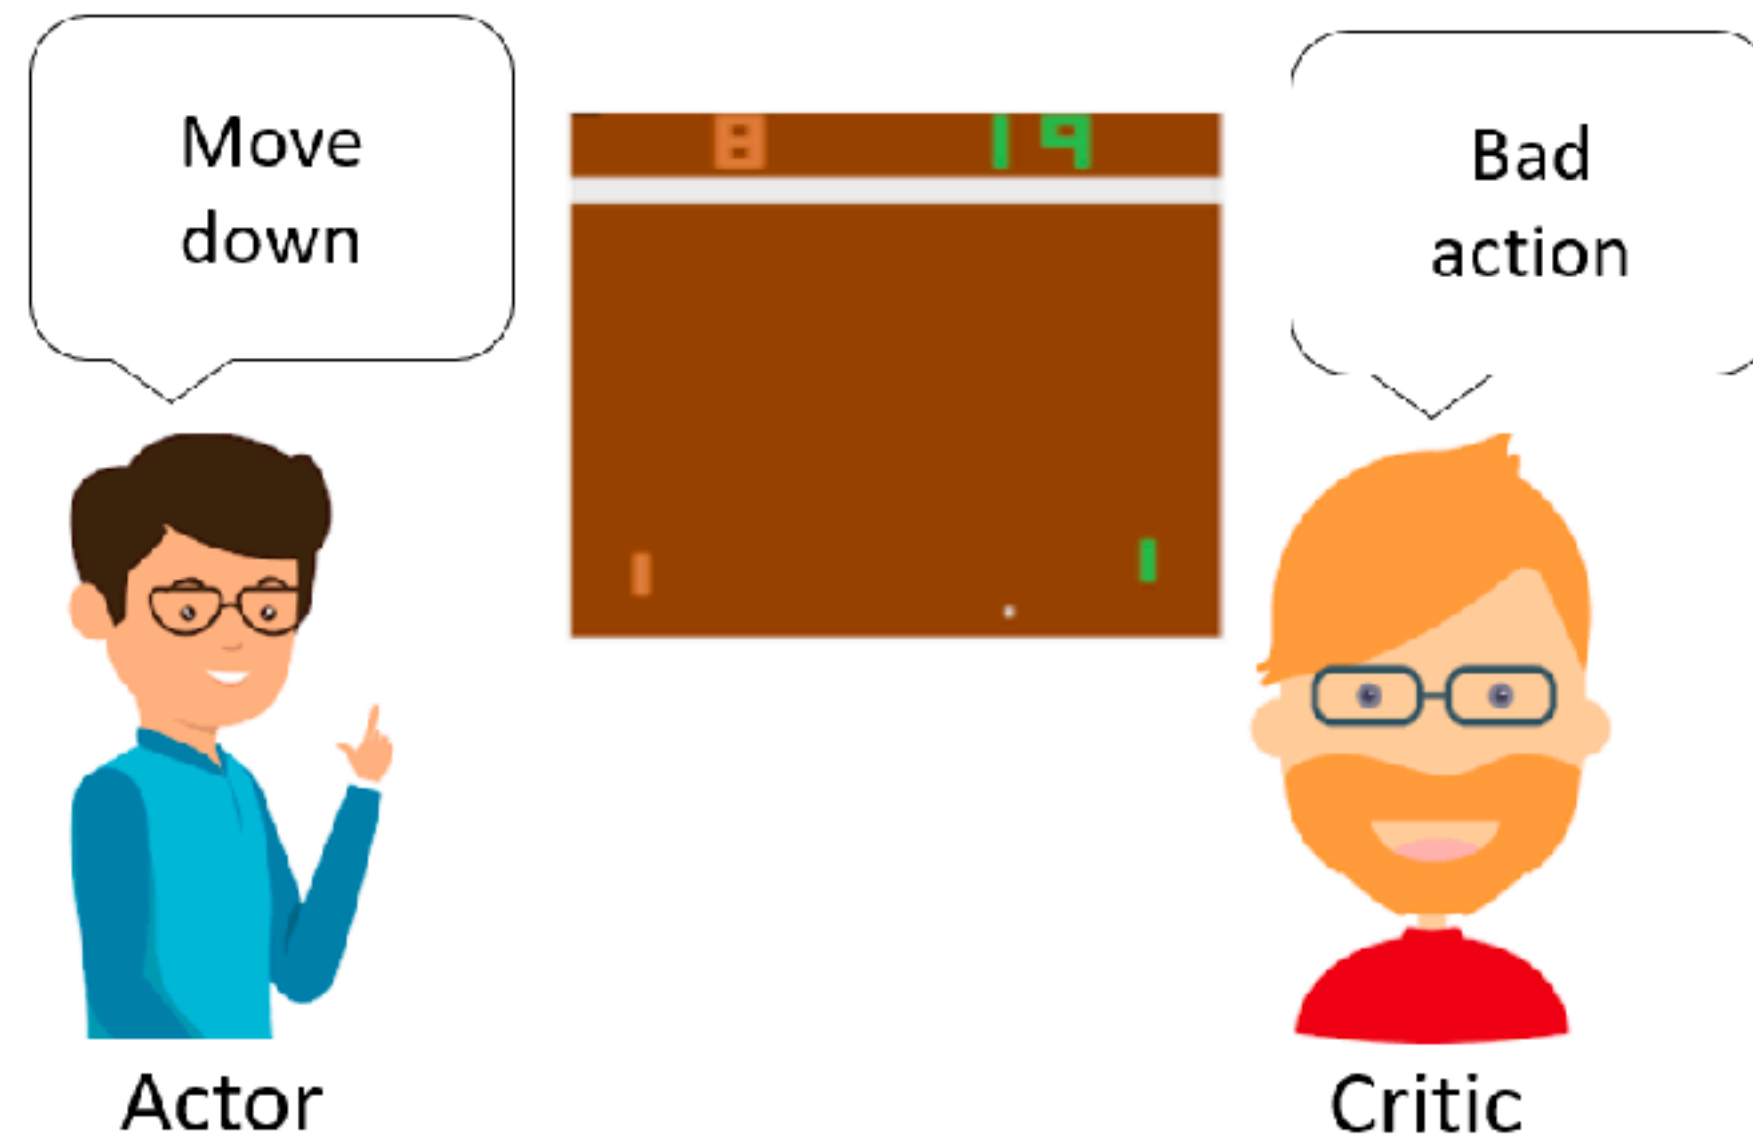
Advantage value



- **Actor**: a PG algorithm that decides on an action to take;
- **Critic**: Q-learning algorithm that critiques the action that the Actor selected, providing feedback on how to adjust. It can take advantage of efficiency tricks in Q-learning, such as memory replay.
- **how better it is to take a specific action than the average general action at the given state?**


A2C

...fit states to actions directly (actor) + evaluate advantage of a new state (critic)




A2C

...fit states to actions directly (actor) + evaluate advantage of a new state (critic) - Model



A2C

...fit states to actions directly (actor) + evaluate advantage of a new state (critic) - Model



```
def OurModel(input_shape, action_space, lr):  
    X_input = Input(input_shape)  
  
    X = Flatten(input_shape=input_shape)(X_input)  
  
    X = Dense(512, activation="elu", kernel_initializer='he_uniform')(X)  
  
    action = Dense(action_space, activation="softmax", kernel_initializer='he_uniform')(X)  
    value = Dense(1, kernel_initializer='he_uniform')(X)  
  
    Actor = Model(inputs = X_input, outputs = action)  
    Actor.compile(loss='categorical_crossentropy', optimizer=RMSprop(lr=lr))  
  
    Critic = Model(inputs = X_input, outputs = value)  
    Critic.compile(loss='mse', optimizer=RMSprop(lr=lr))  
  
    return Actor, Critic
```


A2C

...fit states to actions directly (actor) + evaluate advantage of a new state (critic) - Training

```
def replay(self):  
    # reshape memory to appropriate shape for training  
    states = np.vstack(self.states)  
    actions = np.vstack(self.actions)  
  
    # Compute discounted rewards  
    discounted_r = self.discount_rewards(self.rewards)  
  
    # Get Critic network predictions  
    values = self.Critic.predict(states)[:, 0]  
    # Compute advantages  
    advantages = discounted_r - values  
    # training Actor and Critic networks  
    self.Actor.fit(states, actions, sample_weight=advantages, epochs=1, verbose=0)  
    self.Critic.fit(states, discounted_r, epochs=1, verbose=0)  
    # reset training memory  
    self.states, self.actions, self.rewards = [], [], []
```

A2C

...fit states to actions directly (actor) + evaluate advantage of a new state (critic) - Ray config

```
1  # Run with:
2  # rllib train file cartpole_a2c.py \
3  #     --stop={'timesteps_total': 50000, 'episode_reward_mean': 200}"
4  from ray.rllib.algorithms.a2c import A2CConfig
5
6
7  config = (
8      A2CConfig()
9      .environment("CartPole-v1")
10     .training(lr=0.001, train_batch_size=20)
11     .framework("tf")
12     .rollouts(num_rollout_workers=0)
13 )
```


A2C

Summary

- **A2C learns the policy π directly** (policy head) + accounts for **advantage of a specific action at the given state** (value head)
- **Converge faster** than DQN and PG
- **More stable** than DQN and PG
- **Effective in large action spaces** or using continuous actions.
- Single-threaded (

Section 3

Asynchronous Advantage Actor-Critic (A3C)

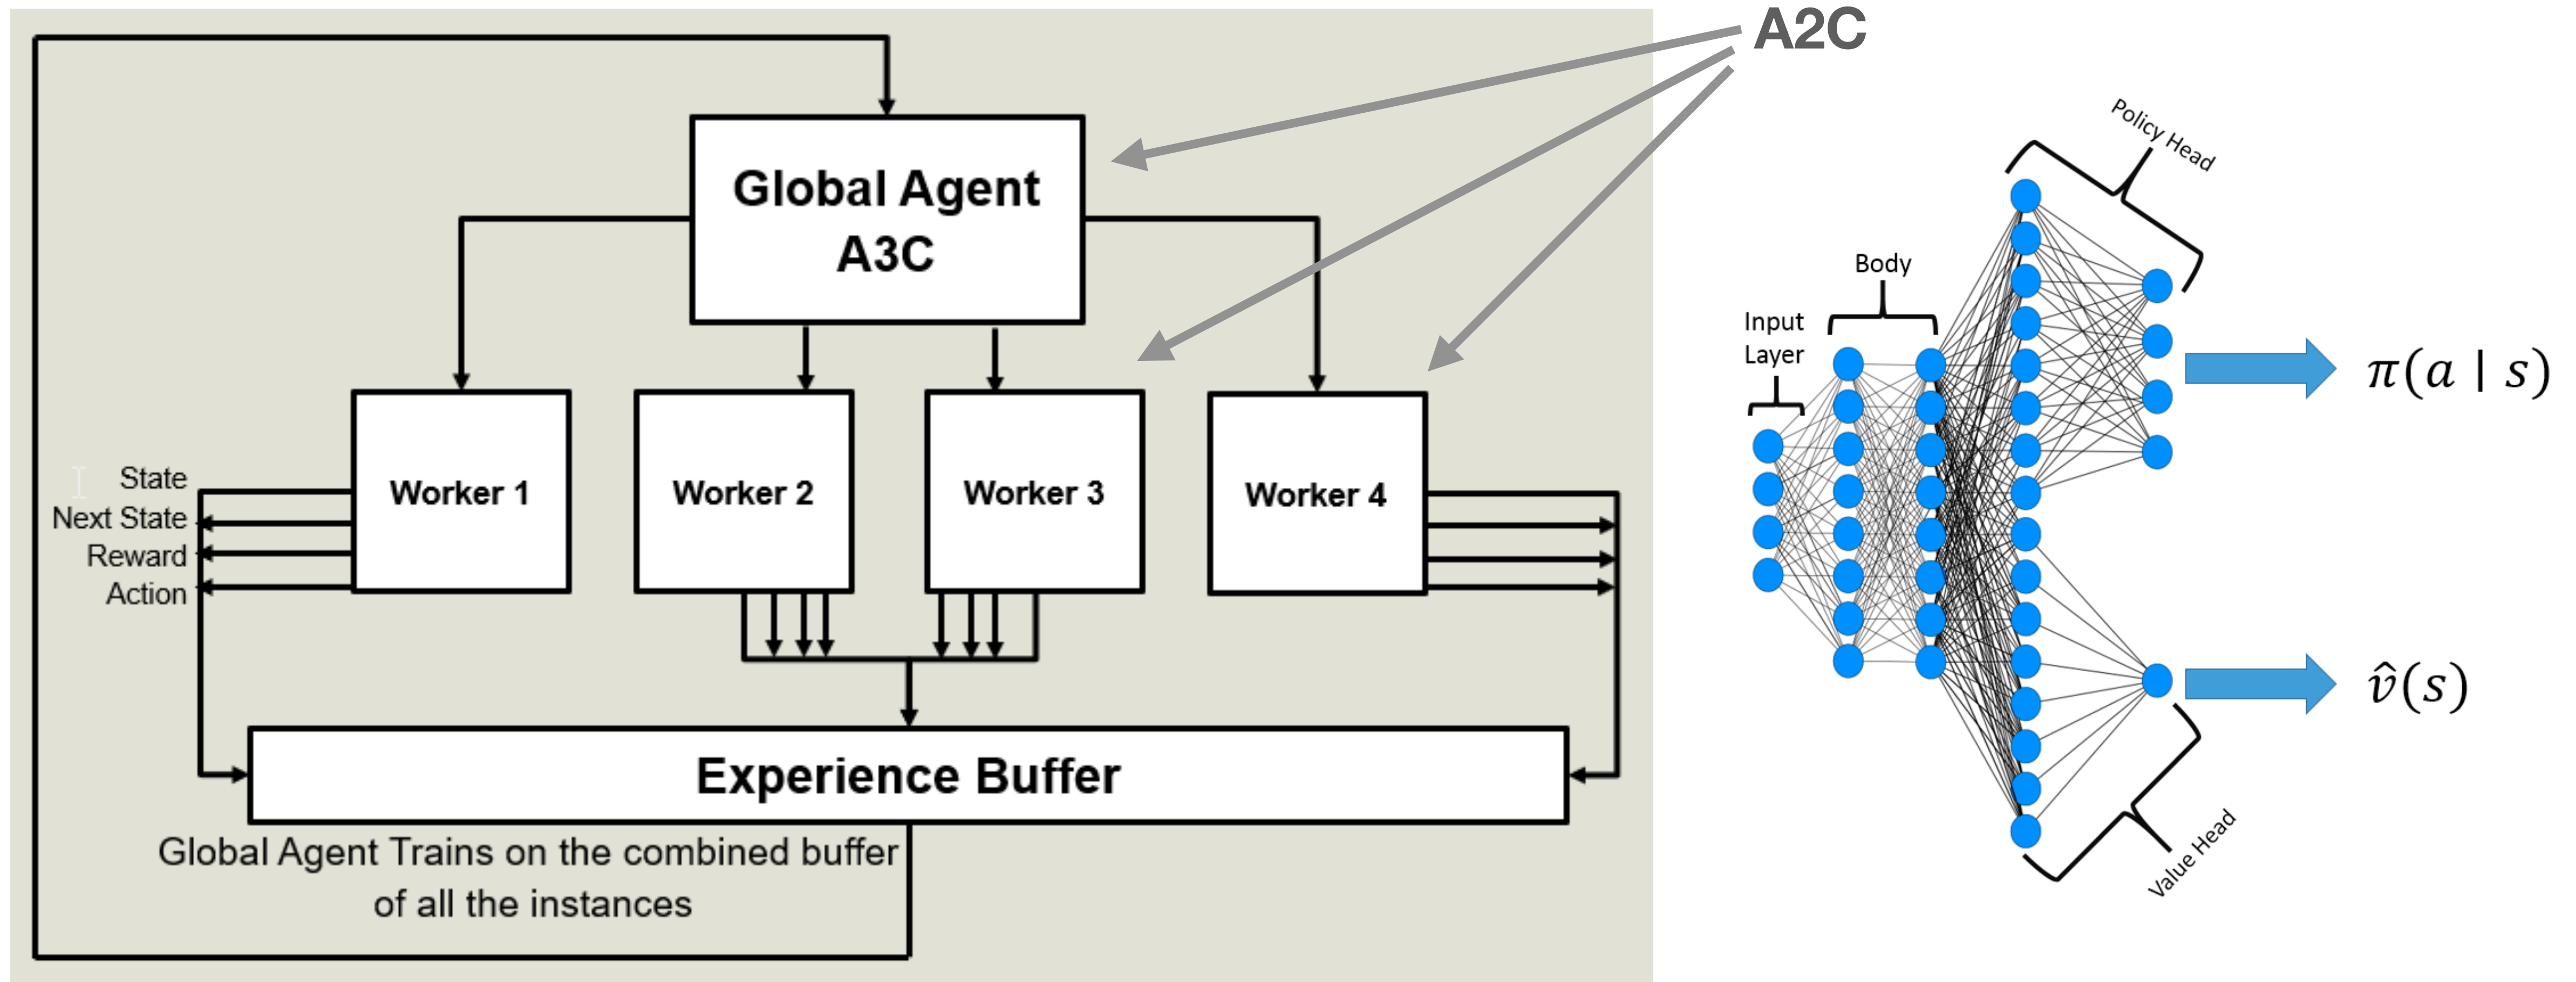
A3C

What's in the name?

- **Asynchronous**: multiple worker agents are trained in parallel, each with their environment.
 - => **train faster** as more workers are training in parallel
 - => **diverse training experience** as each worker's experience is independent
- **Advantage**: a metric to judge how good its actions were and how they turned out.
 - => **measure the advantage of an action** at time step t , following the policy π
 - => **focus** on where the network's predictions were lacking
- **Actor-Critic**: the algorithm's architecture shares layers between the policy and value function.

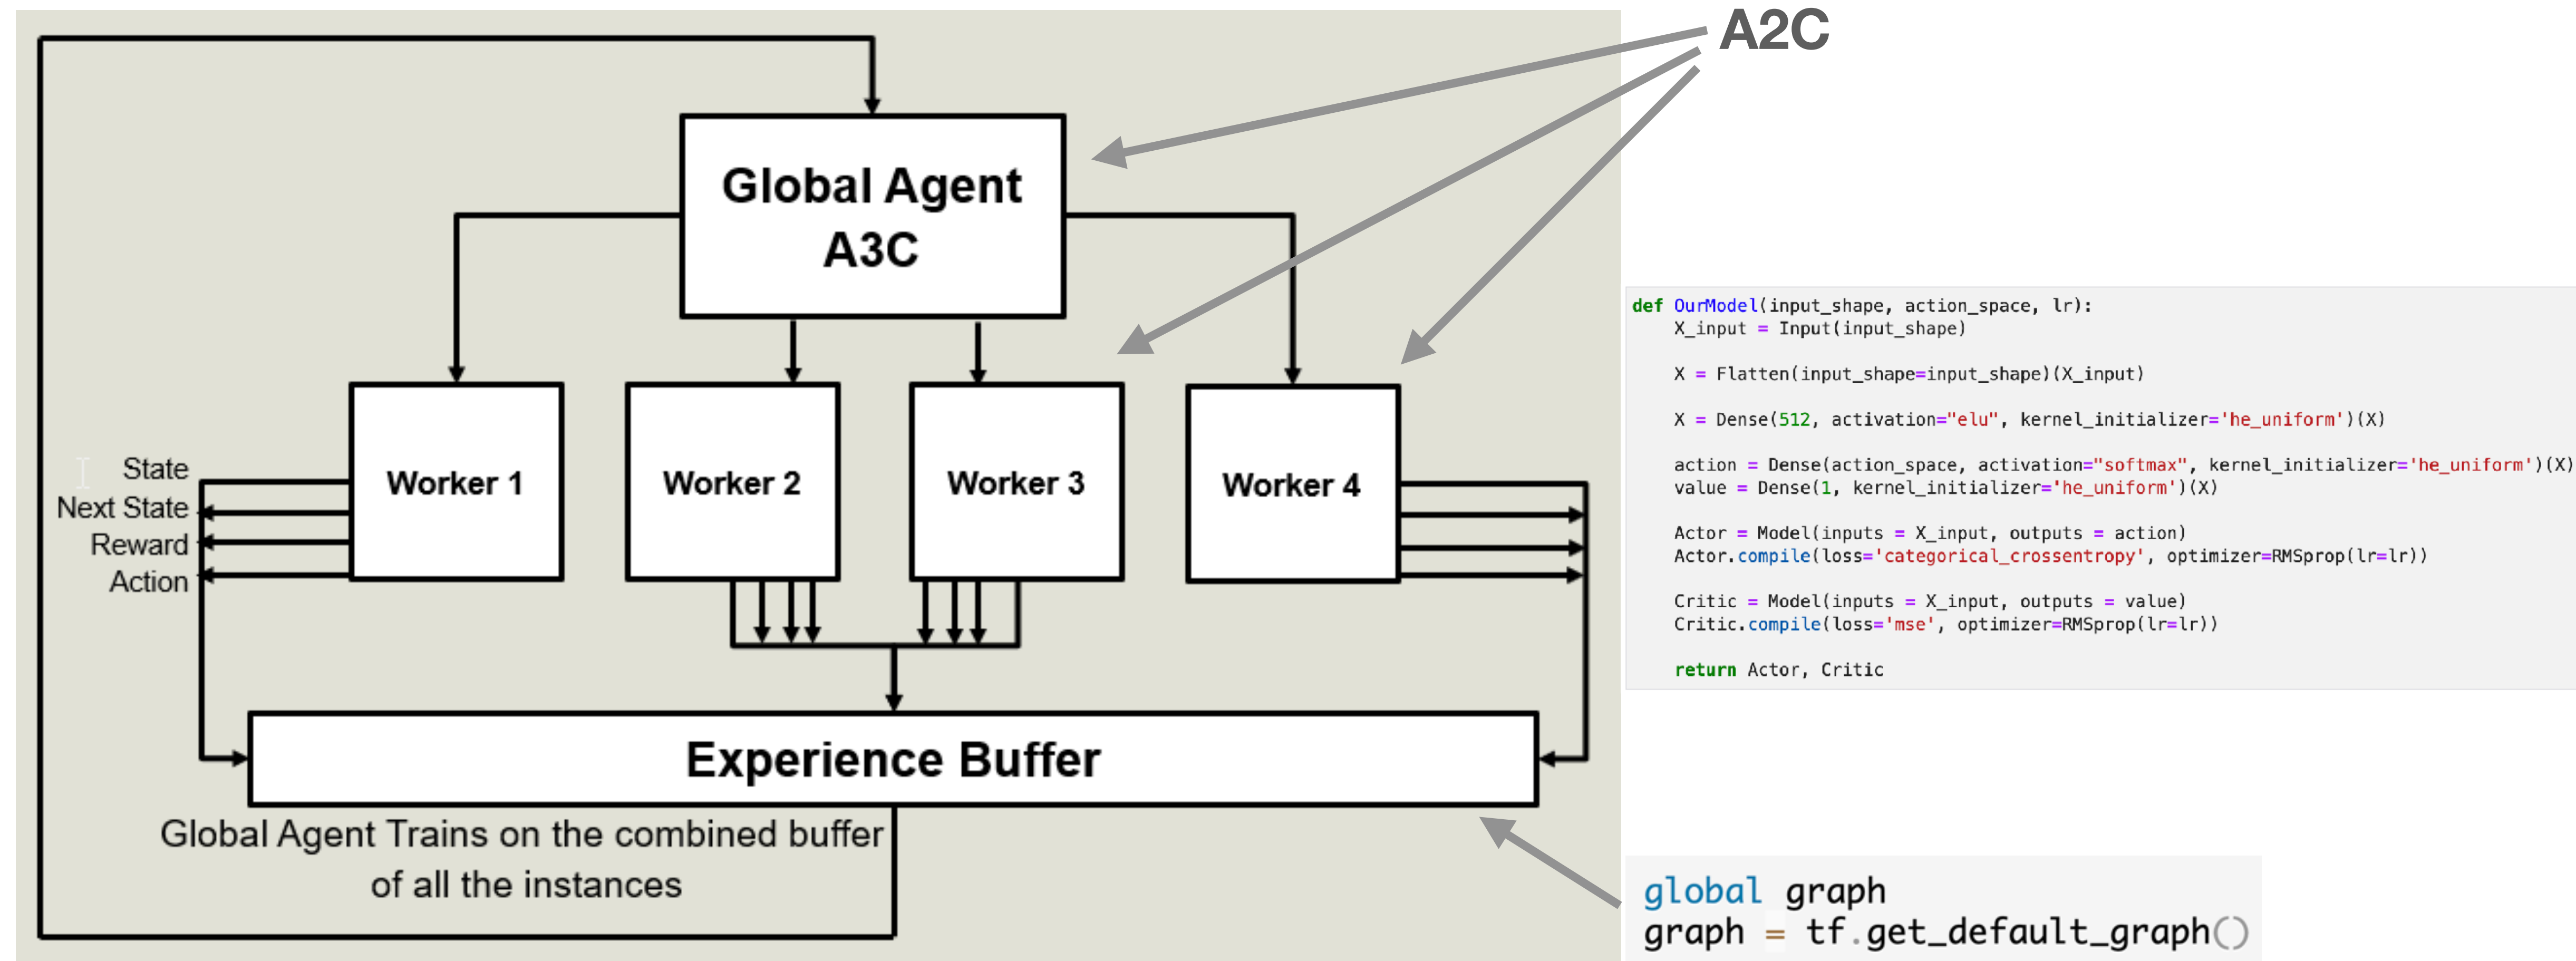
A3C

...is the asynchronous training technology + A2C neural nets



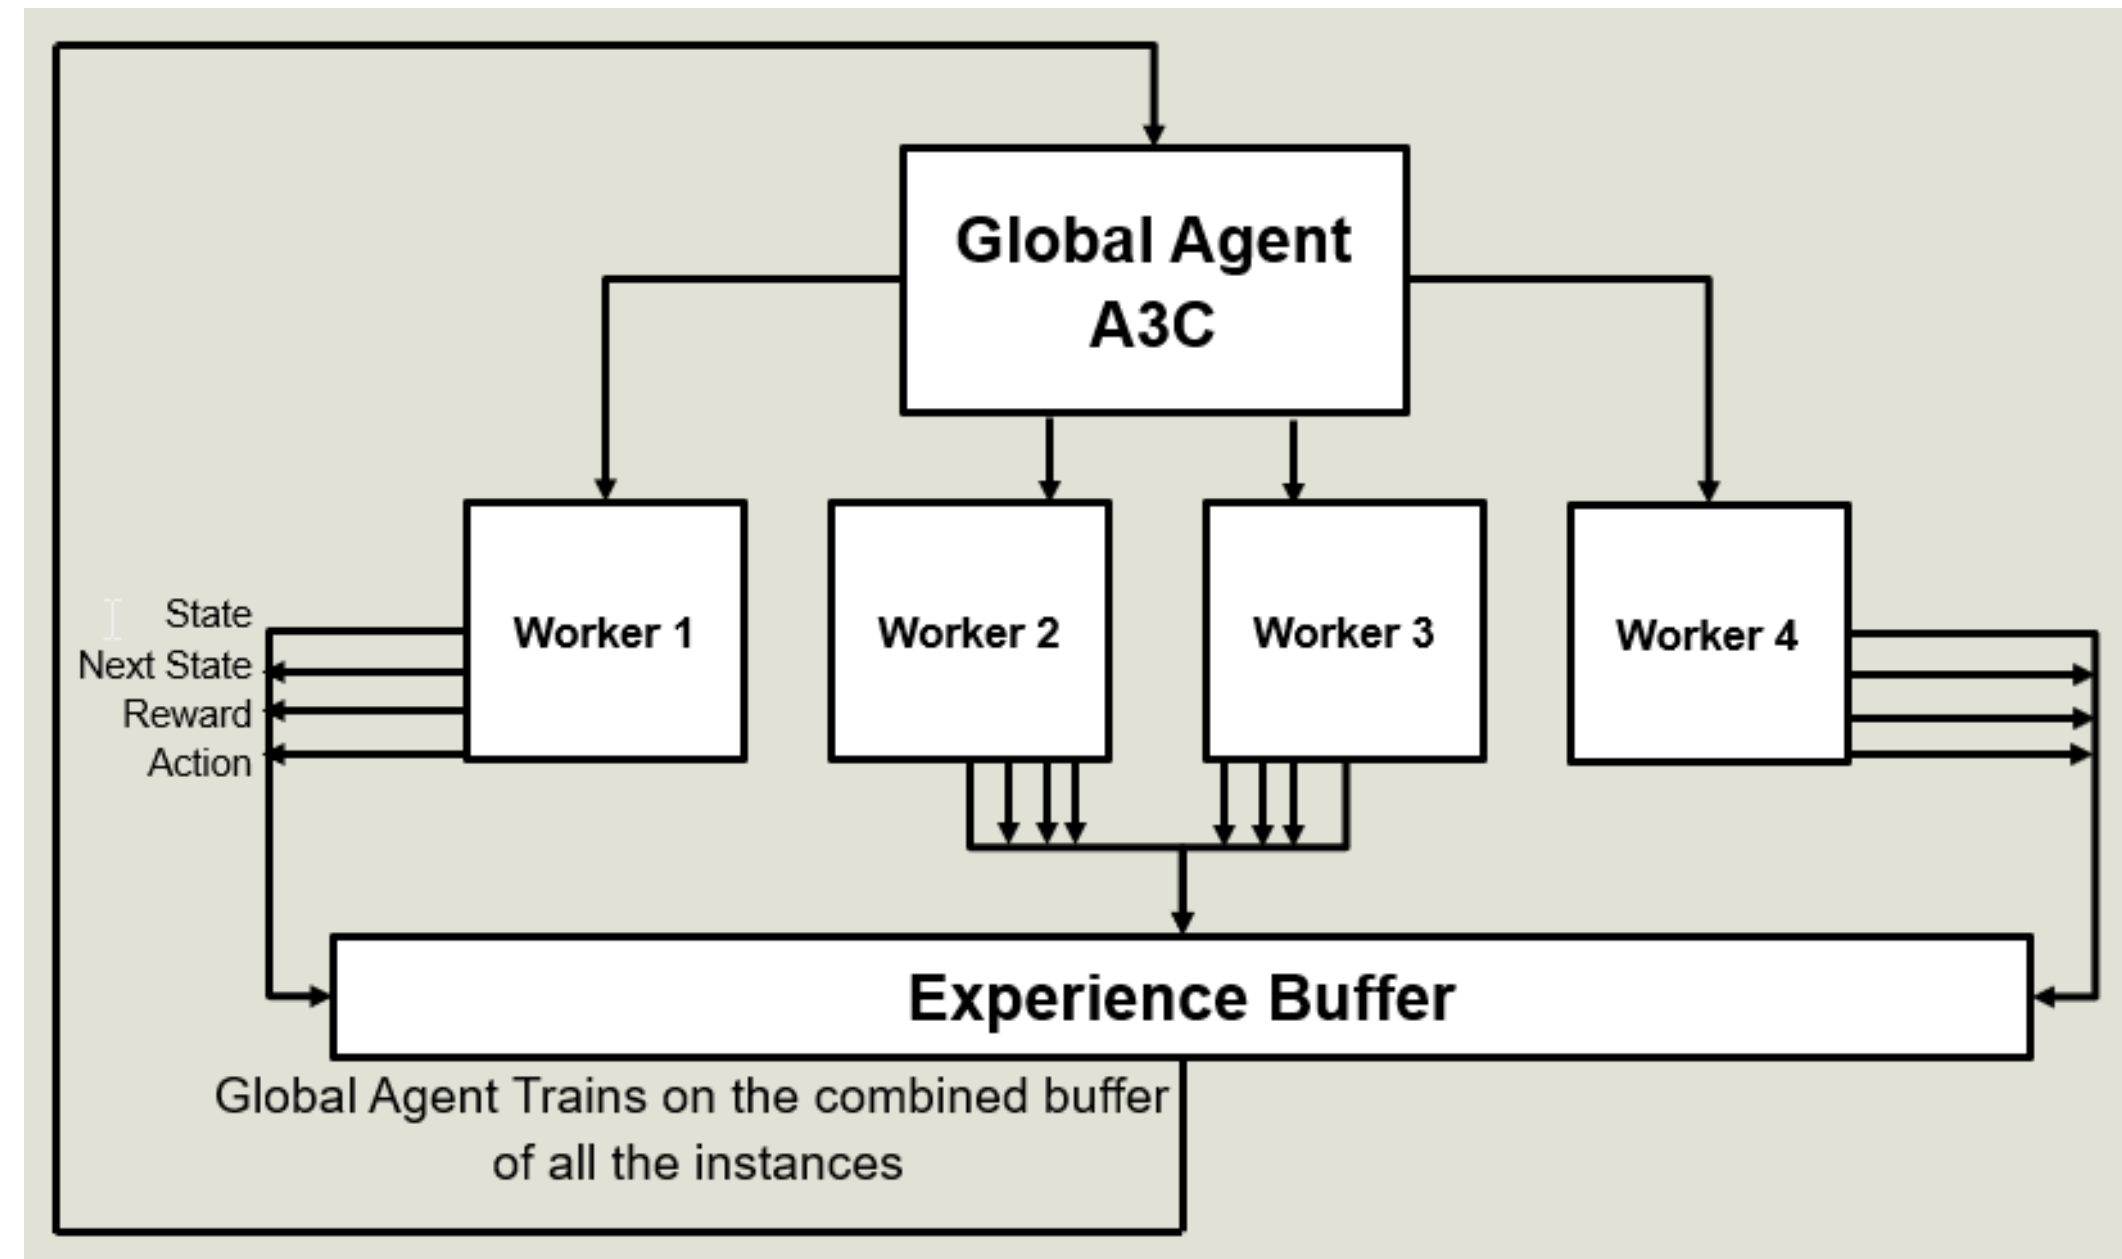
A3C

...is the asynchronous training technology + A2C neural nets



A3C

How it works?



- Each **Worker** trains independently in its own environment
- Each **Worker** at the end of each episode writes its local Experience Buffer to a global **Experience Buffer**
- The **Global Agent** then trains on the global **Experience Buffer**
- The workers then copy the weights from **Global Agent**
- ...repeat until **Global Agent** converges

A3C

How it works?

```
def train(self, n_threads):
    self.env.close()
    # Instantiate one environment per thread
    envs = [gym.make(self.env_name) for i in range(n_threads)]

    # Create threads
    threads = [threading.Thread(
        target=self.train_threading,
        daemon=True,
        args=(self,
              envs[i],
              i)) for i in range(n_threads)]

    for t in threads:
        time.sleep(2)
        t.start()
```

**Training
(multi-threaded)**

Experience Buffer

```
global graph
graph = tf.get_default_graph()
```

**One worker's environment
(single thread)**

```
def train_threading(self, agent, env, thread):
    global graph
    with graph.as_default():
        while self.episode < self.EPISODES:
            # Reset episode
            score, done, SAVING = 0, False, ''
            state = self.reset(env)
            # Instantiate or reset games memory
            states, actions, rewards = [], [], []
            while not done:
                action = agent.act(state)
                next_state, reward, done, _ = self.step(action, env, state)

                states.append(state)
                action_onehot = np.zeros([self.action_size])
                action_onehot[action] = 1
                actions.append(action_onehot)
                rewards.append(reward)

                score += reward
                state = next_state

            self.lock.acquire()
            self.replay(states, actions, rewards)
            self.lock.release()
```

```
def replay(self, states, actions, rewards):
    # reshape memory to appropriate shape for training
    states = np.vstack(states)
    actions = np.vstack(actions)

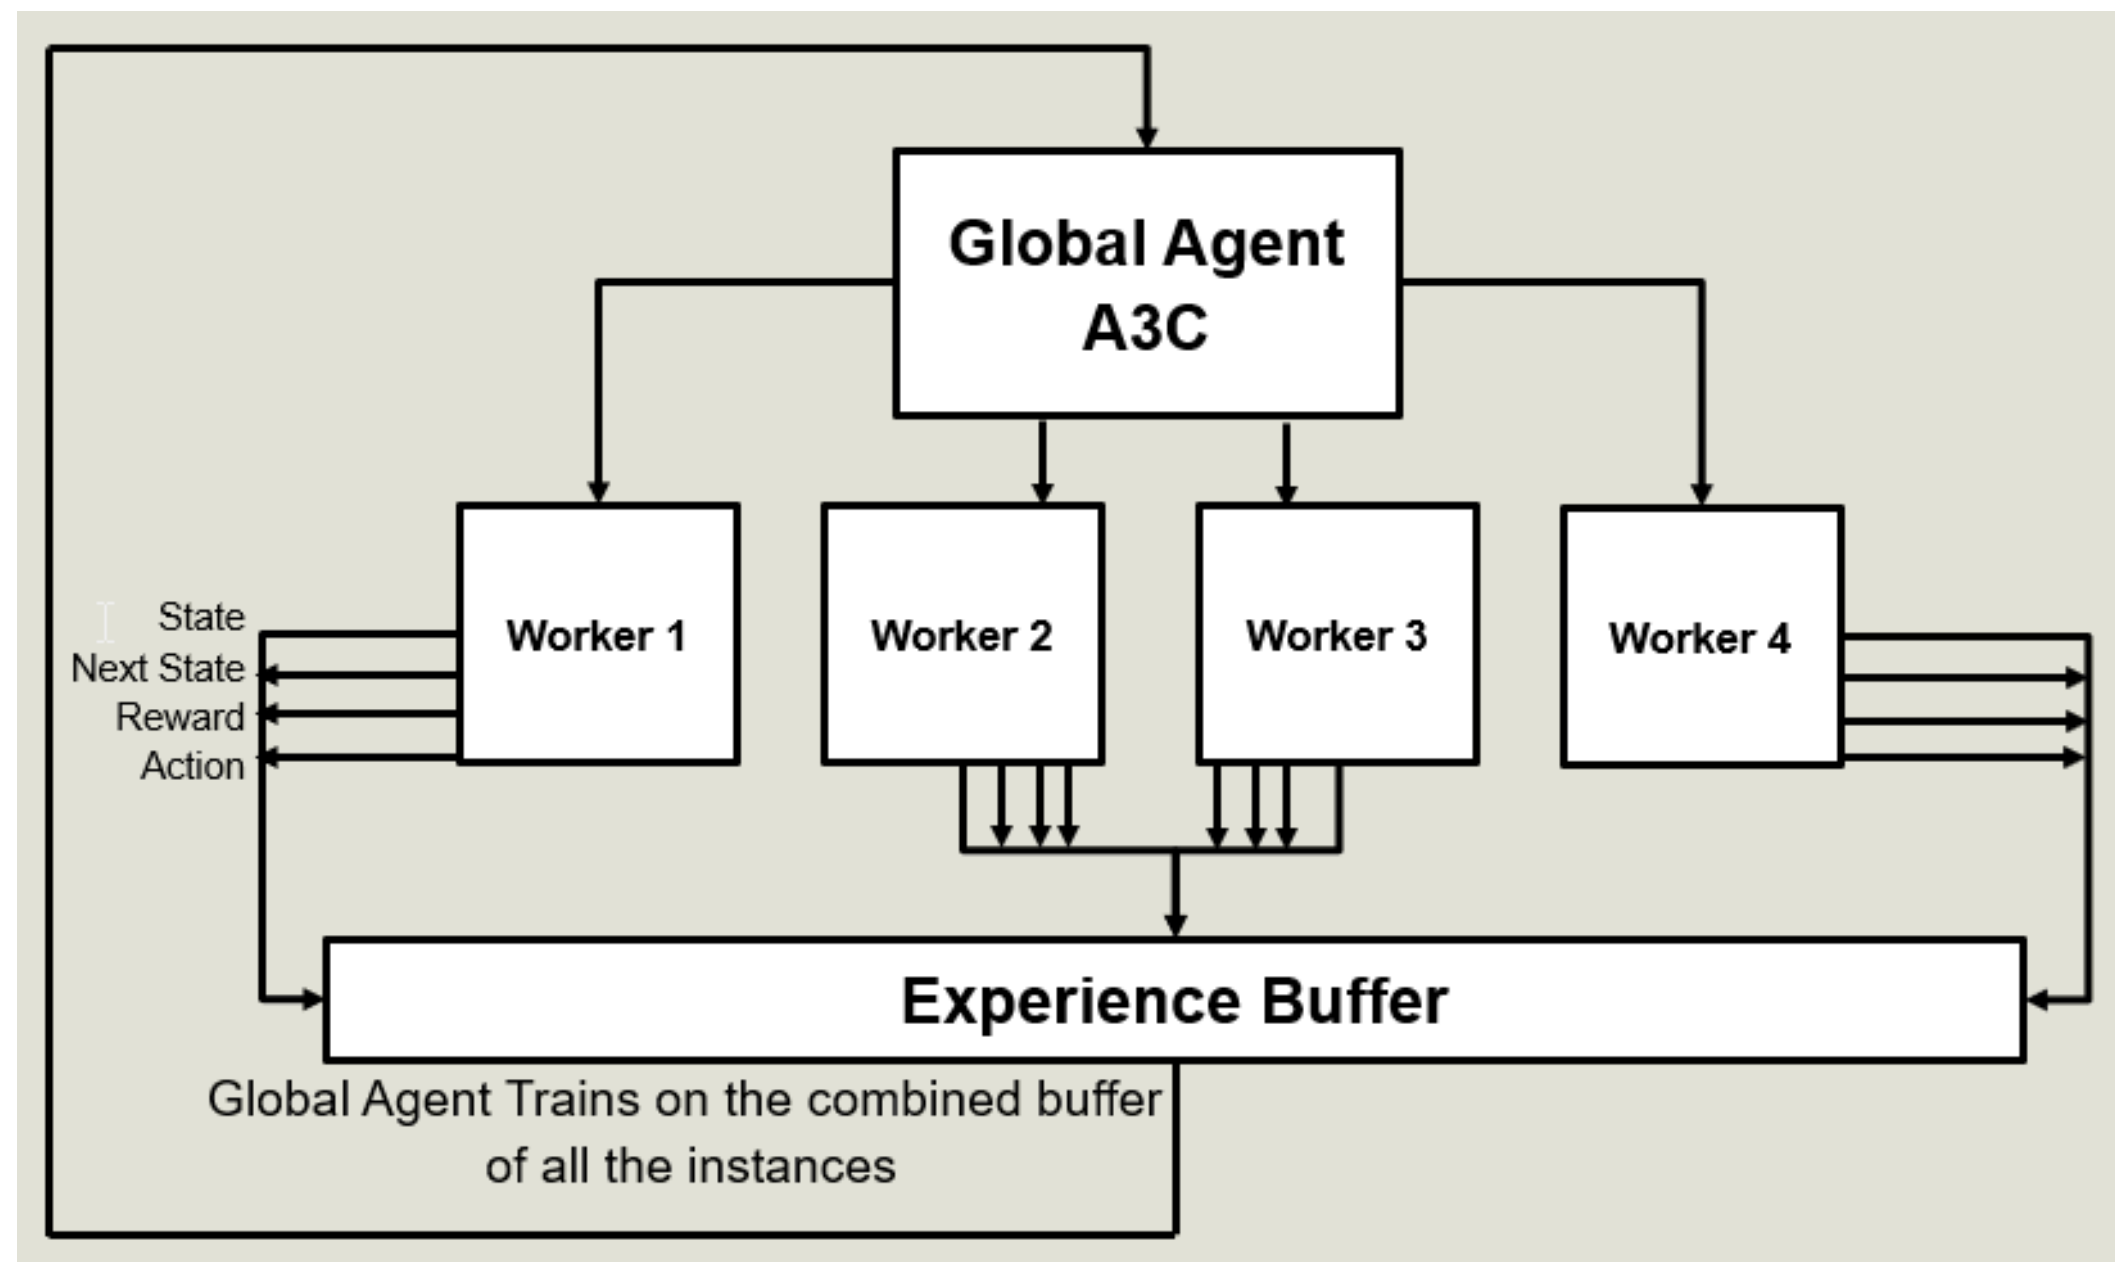
    # Compute discounted rewards
    discounted_r = self.discount_rewards(rewards)

    # Get Critic network predictions
    value = self.Critic.predict(states[:, 0])
    # Compute advantages
    advantages = discounted_r - value
    # training Actor and Critic networks
    self.Actor.fit(states, actions, sample_weight=advantages, epochs=1, verbose=0)
    self.Critic.fit(states, discounted_r, epochs=1, verbose=0)
```

**Training on
Experience Buffer**

A3C

What is the point?



- Faster training since workers running in parallel.
- Distributed / federated training

A3C

Ray config (.yaml & .py)

```
3  pong-a3c:
4      env: PongDeterministic-v4
5      run: A3C
6      config:
7          # Works for both torch and tf.
8          framework: tf
9          num_workers: 16
10         rollout_fragment_length: 20
11         vf_loss_coeff: 0.5
12         entropy_coeff: 0.01
13         gamma: 0.99
14         grad_clip: 40.0
15         lambda: 1.0
16         lr: 0.0001
17         observation_filter: NoFilter
18         preprocessor_pref: rllib
19         model:
20             use_lstm: true
21             conv_activation: elu
22             dim: 42
23             grayscale: true
24             zero_mean: false
25             # Reduced channel depth and kernel size from default
26             conv_filters: [
27                 [32, [3, 3], 2],
28                 [32, [3, 3], 2],
29                 [32, [3, 3], 2],
30                 [32, [3, 3], 2],
31             ]
```

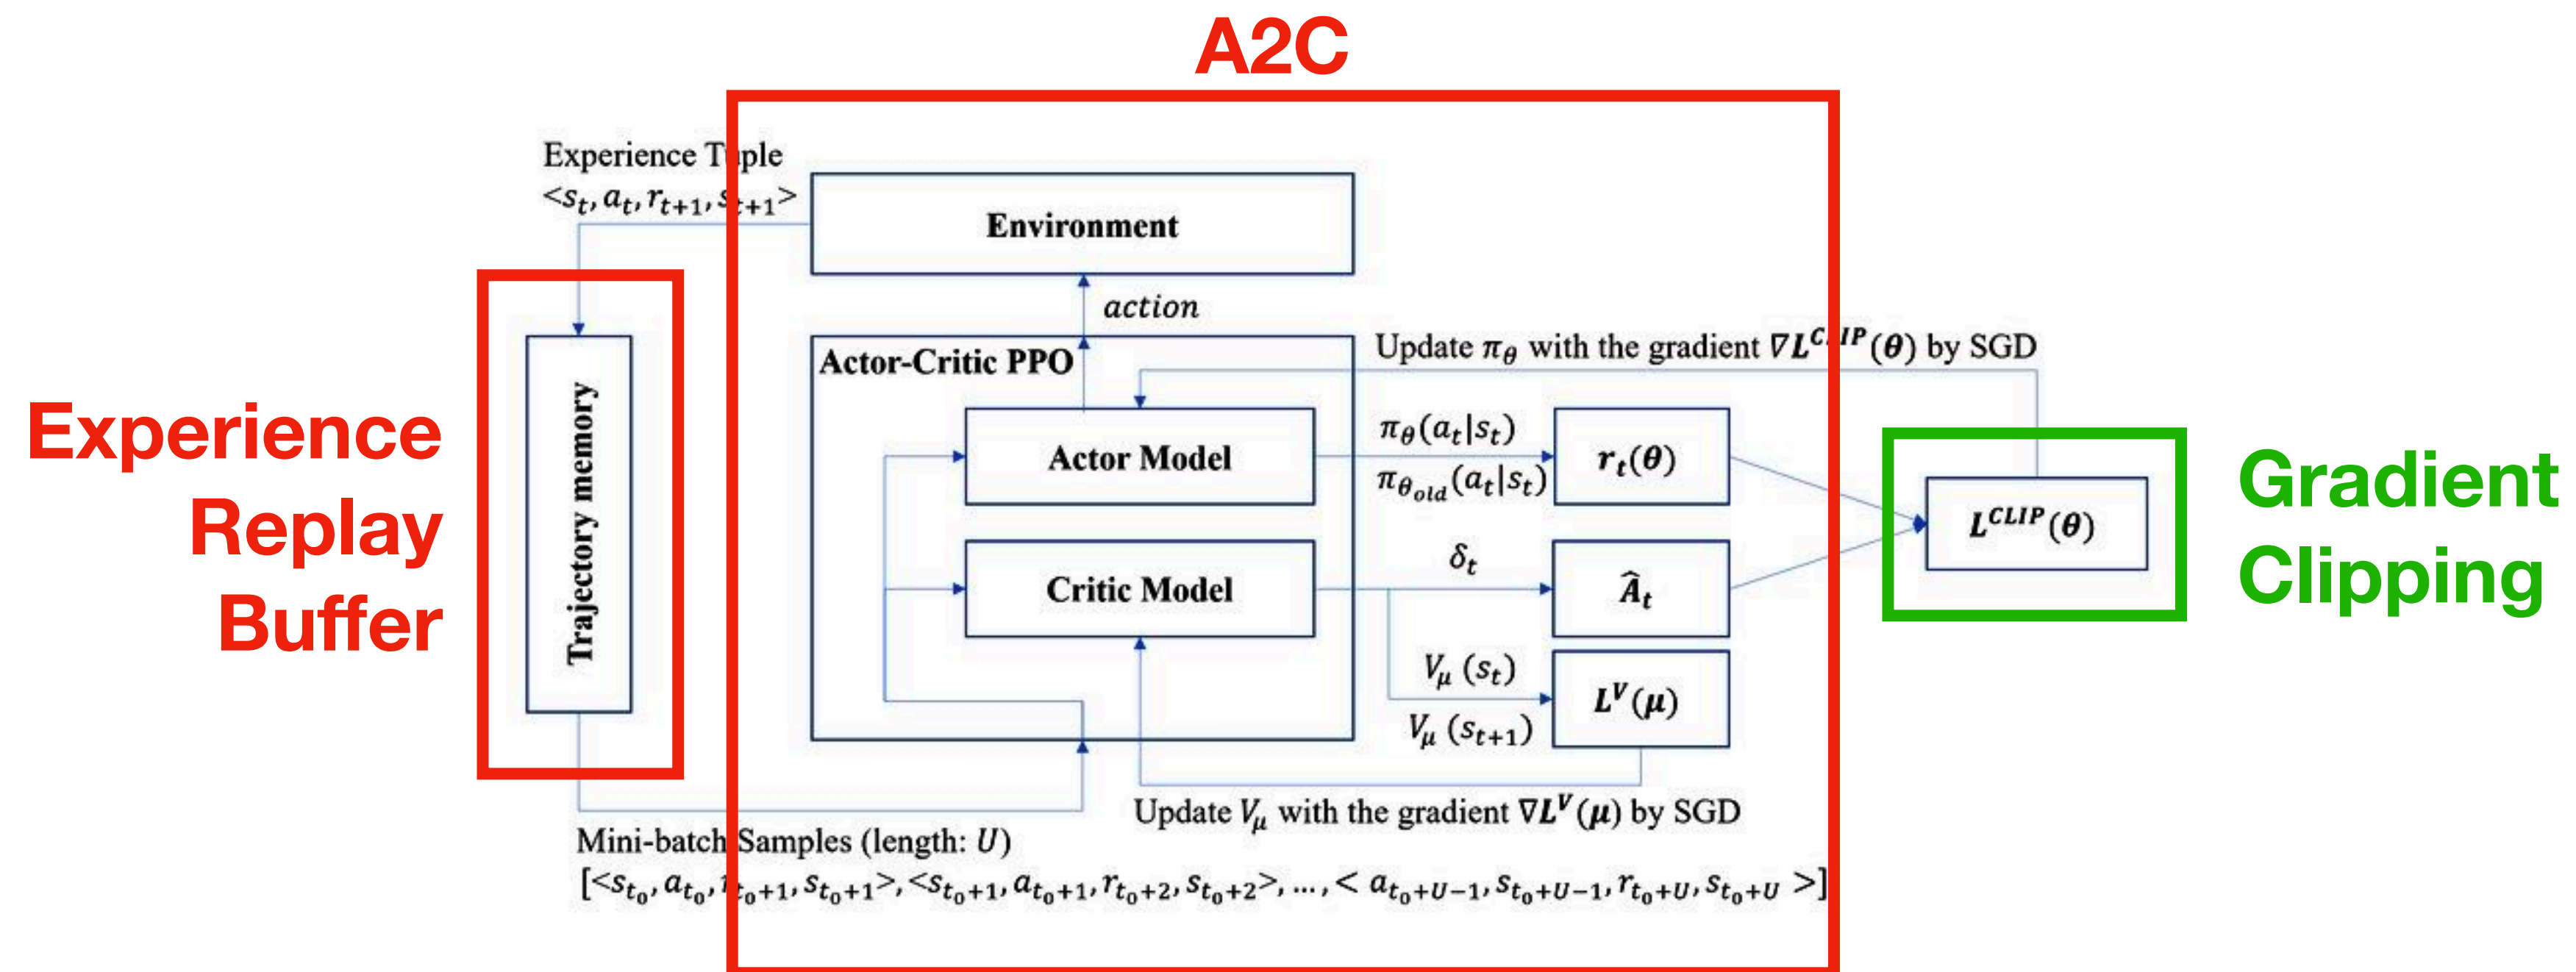
```
1  # Run with:
2  # rllib train file cartpole_a3c.py \
3  #     --stop={'timesteps_total': 20000, 'episode_reward_mean': 150}"
4  from ray.rllib.algorithms.a3c import A3CConfig
5
6
7  config = (
8      A3CConfig()
9      .training(gamma=0.95)
10     .environment("CartPole-v1")
11     .framework("tf")
12     .rollouts(num_rollout_workers=0)
13 )
```

Section 4

Proximal Policy Optimisation (PPO)

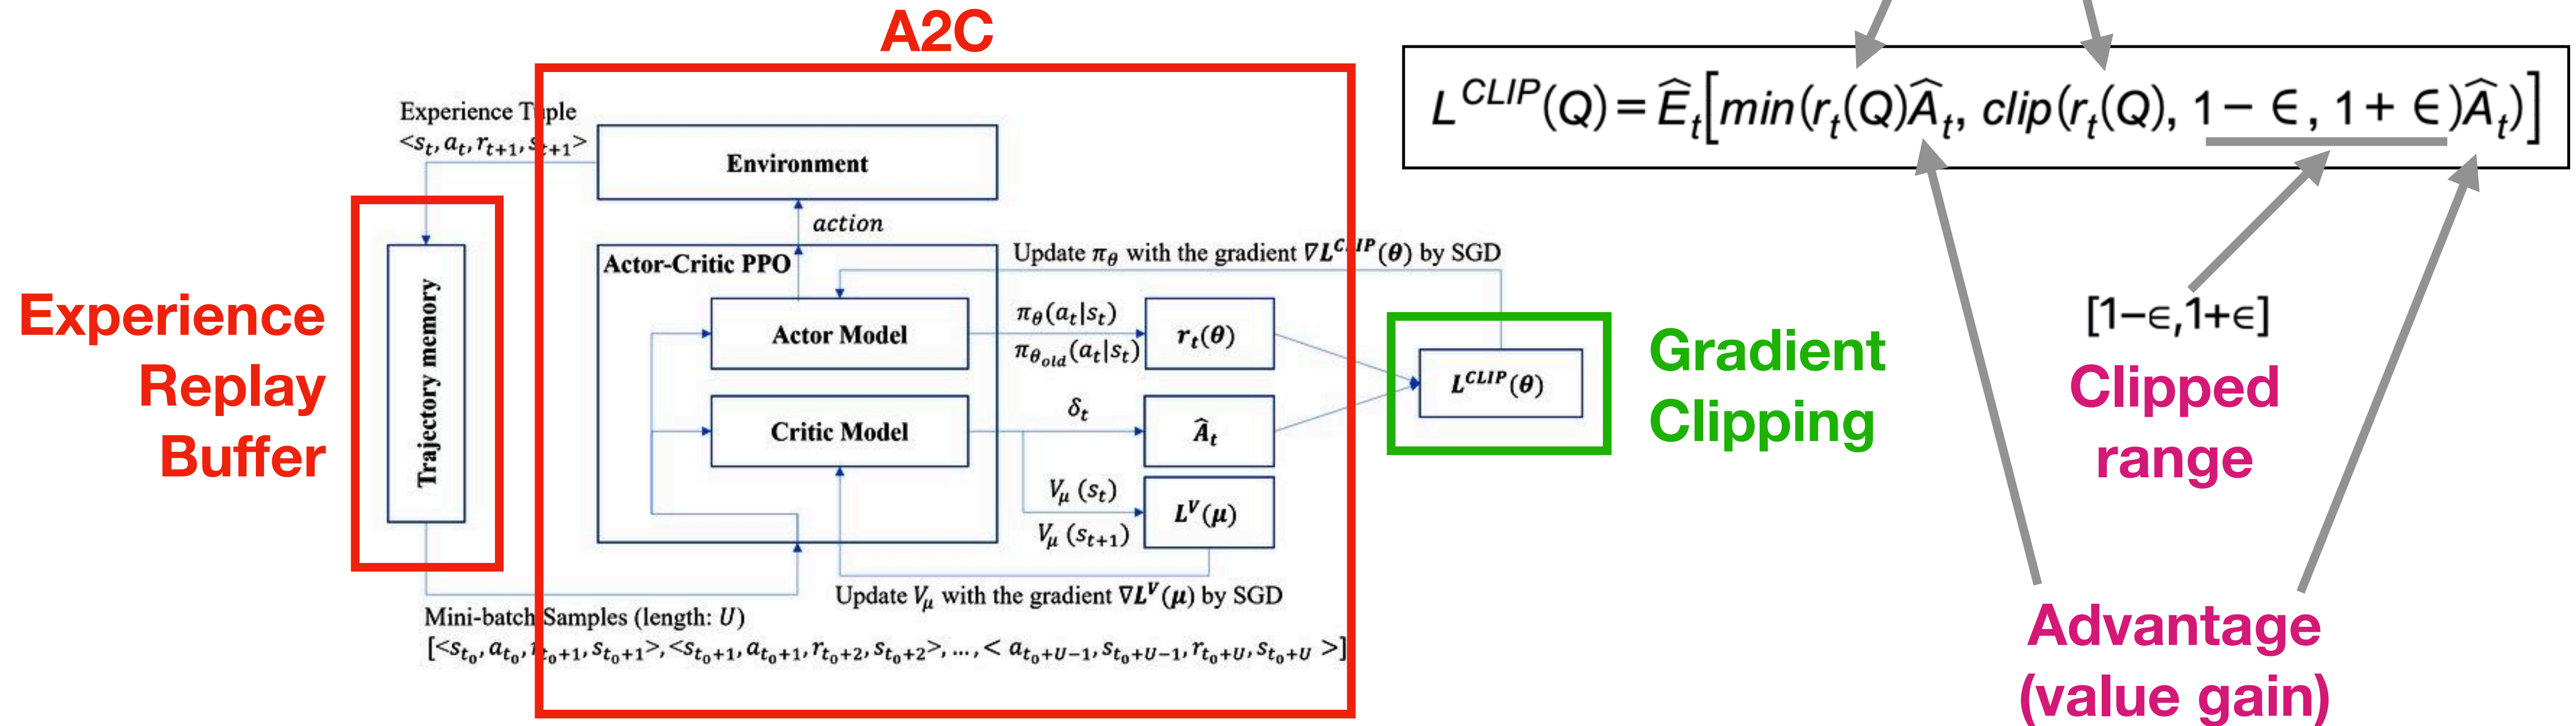
PPO

= A2C + policy gradient clipping



PPO

= A2C + policy gradient clipping



PPO

Model architecture

```
def OurModel(input_shape, action_space, lr):
    X_input = Input(input_shape)

    X = Flatten(input_shape=input_shape)(X_input)

    X = Dense(512, activation="elu", kernel_initializer='he_uniform')(X)

    action = Dense(action_space, activation="softmax", kernel_initializer='he_uniform')(X)
    value = Dense(1, activation='linear', kernel_initializer='he_uniform')(X)

    def ppo_loss(y_true, y_pred):
        # Defined in https://arxiv.org/abs/1707.06347
        advantages, prediction_picks, actions = \
            y_true[:, :1], y_true[:, 1:1+action_space], y_true[:, 1+action_space:]
        LOSS_CLIPPING = 0.2
        ENTROPY_LOSS = 5e-3

        prob = y_pred * actions
        old_prob = actions * prediction_picks
        r = prob/(old_prob + 1e-10)
        p1 = r * advantages
        p2 = K.clip(r, min_value=1 - LOSS_CLIPPING, max_value=1 + LOSS_CLIPPING) * advantages
        loss = -K.mean(K.minimum(p1, p2) + ENTROPY_LOSS * -(prob * K.log(prob + 1e-10)))

        return loss

    Actor = Model(inputs = X_input, outputs = action)
    Actor.compile(loss=ppo_loss, optimizer=RMSprop(lr=lr))

    Critic = Model(inputs = X_input, outputs = value)
    Critic.compile(loss='mse', optimizer=RMSprop(lr=lr))

    return Actor, Critic
```

$$L^{CLIP}(Q) = \hat{E}_t \left[\min(r_t(Q) \hat{A}_t, \text{clip}(r_t(Q), 1 - \epsilon, 1 + \epsilon) \hat{A}_t) \right]$$

PPO

Ray config (.yaml)

```
3 # $ python train.py -f tuned_configs/pong-ppo.yaml
4 #
5 pong-ppo:
6   env: PongNoFrameskip-v4
7   run: PP0
8   config:
9     # Works for both torch and tf.
10    framework: tf
11    lambda: 0.95
12    kl_coeff: 0.5
13    clip_rewards: True
14    clip_param: 0.1
15    vf_clip_param: 10.0
16    entropy_coeff: 0.01
17    train_batch_size: 5000
18    rollout_fragment_length: 20
19    sgd_minibatch_size: 500
20    num_sgd_iter: 10
21    num_workers: 32
22    num_envs_per_worker: 5
23    batch_mode: truncate_episodes
24    observation_filter: NoFilter
25    num_gpus: 1
26    model:
27      dim: 42
28      vf_share_layers: true
```


Homework

Assault + Gym Env + Ray RLlib

- **Train & tune A2C on Assault**

Tuned examples: https://github.com/ray-project/ray/tree/master/rllib/tuned_examples/a2c

- **Train & tune A3C on Assault**

Tuned examples: https://github.com/ray-project/ray/tree/master/rllib/tuned_examples/a3c

- **Train & tune PPO on Assault**

Tuned examples: https://github.com/ray-project/ray/tree/master/rllib/tuned_examples/ppo